

```
#Name: Abhishek N. Singh  
#Course: Matrix Decomposition for Data Analysis  
#Entry No.: 322079  
#
```

```
from util import *  
#from utils import * #Note this is not needed  
import numpy as np  
from scipy.sparse import coo_matrix  
import matplotlib.pyplot as plt  
import time
```

```
#!pip install utils #This is not needed
```

```
Collecting utils  
  Downloading https://files.pythonhosted.org/packages/55/e6/c2d2b2703e7debc8b501caaee0e6  
Installing collected packages: utils  
Successfully installed utils-1.0.1
```



```
#Task 1: Matrix completion for movie ratings  
#Download the data and code stub files from course's Moodle page. Among the files, there are  
#that we use in this task: ratings.csv, ratings_validation.csv, and movies.txt. This data is  
#movie ratings data from GroupLens Research.  
#It contains ratings (in 0 to 5 stars scale) of 9724 movies by  
#610 users. Most of these ratings are in ratings.csv. The last rating (by timestamp) of every  
#in the ratings_validation.csv file. We use this to evaluate how good our prediction of the r  
#Do not use the validation data to train your model! The final file, movies.txt, contains the  
#movies in the same order as they appear in the data (i.e. the first movie name corresponds t  
#1, etc.).  
#First, load the data and normalize the ratings so that they're first in scale from 1 to 6 st  
#them by 6 so that the final ratings are in interval (0, 1]. We will use 0 to denote a missir  
#can use standard sparse matrix storage formats.  
#The code stub contains an implementation of an SGD algorithm. For this and the next task you  
#this implementation so that the results are comparable. The code can be slow, so reserve enc  
#complete these tasks!  
#Now, get some basic idea how SGD algorithm works: run the provided SGD algorithm with this c  
#using learning rate 0.001 and 100 epochs and no regularization, and observe how the error be
```

```

#epochs. Plot errors over epochs. Has the algorithm converged in 100 epochs? Argue!
#The algorithm does not do early stopping if the change between two consecutive epochs is sma
#Would this make sense for an SGD algorithm? Why or why not?
#You can also print the top-10 and bottom-10 movies per factor. Can you identify any movies c
#patterns? Include some examples in your report (but not all).
#In the previous run you used a fixed learning rate. A common heuristic to set dynamic learni
#the so-called bold driver heuristic. In bold driver heuristic, we increase the learning rate
#decreases over epochs. If the error increases, we drop the learning rate back to the base le
#provided function implements the bold driver heuristic with parameter bold_driver which is n
#with the learning rate when error decreases.3 Run SGD again, this time with learning rate 0.
#driver multiplier of 2. How does the behaviour change? Do you get better results? How is the
#Now, try around different step sizes, bold driver multipliers, and epoch numbers. Try only f
#values so that this part doesn't take too long; we will return to this topic in the next ass
#converge to (essentially) equally good training results with smaller number of epochs? How c
#the validation error?

#Next, we will work with biases. Start by running SGD with learning rate 0.0001 and bold driv
#2 but computing only the global, user, and movie bias. Hence, the prediction of a rating for
#i is  $\mu + b$ 
#user
# $i + b$ 
#movie
#j
#, where  $\mu$  is the average rating. How are the results? Why? Argue!

```

```

#Now, using the same settings, compute both the factor matrices and the bias terms. Now the p
#rating for element (i, j) becomes
# $\mu + b$ 
#user
# $i + b$ 
#movie
# $j + l_i$ 
#.  $r_j = \mu + b$ 
#user
# $i + b$ 
#movie
# $j + (LRT$ 
# $)_{ij}$ ,
#where  $l_i$ 
#is the ith row of left-hand factor matrix L and  $r_j$ 
#is the jth row or right-hand factor matrix R (here,
#R is transposed). Are the results better? Also print the movies with highest and lowest bias
#any patterns here?
#As the final part, we add an L2-regularizer for all terms. The error we now try to minimize
# $\sum$ 
# $(i,j) \in \Omega$ 
# $(D_{i,j} - \mu - b$ 
#user
# $i - b$ 
#movie
# $j - l_i$ 

```

```

#· rj)
#2 + λf kLk
#2
#F + λf kRk
#2
#F + λbkb
#userk
#2
#2 + λbkb
#moviek
#2
#2
#,,
#where λf and λb are the factor and bias term regularizer constants, respectively (the provic
#more using regularizer coefficient value 0.001 for both factors and biases. How do the resu]
#do the validation error results change? Why? Argue!
#Use the parameters (epoch number, step size, and bold driver multiplier) you found earlier t
#bias terms and regularizers. How are the results now?
#As the final task, compute the SVD of the rating matrix (assuming 0s mean 0s, not missing va
#compare the error in observed entries and the validation error.

```

```

tmp = np.genfromtxt('ratings.csv', delimiter=',', dtype='int32')
#Reading the data

```

```
tmp
```

```

array([[ 1,    1,    4],
       [ 1,    2,    4],
       [ 1,    3,    4],
       ...,
       [ 610, 9722,    3],
       [ 610, 9723,    3],
       [ 610, 9724,    3]], dtype=int32)

```

```

#Since the rating is on the column number 3 which has column index 2, we modify it
#That is normalization is done between value (0,6]
val = np.array(tmp[:,2]+1, dtype='float')/6

```

```
val
```

```

array([0.83333333, 0.83333333, 0.83333333, ..., 0.66666667, 0.66666667,
       0.66666667])

```

```

#coo_matrix((data, (i, j)), [shape=(M, N)])
#to construct from three arrays:
#data[:] the entries of the matrix, in any order
#i[:] the row indices of the matrix entries

```

```
#j[:] the column indices of the matrix entries
```

```
#Where A[i[k], j[k]] = data[k]. When shape is not specified, it is inferred from the index array
#So now the modified data is:
data = coo_matrix( (val, (tmp[:,0]-1, tmp[:,1]-1)) )
# data is 1-indexed, so we need to reduce 1 to get it 0-indexed
```

```
data
```

```
<610x9724 sparse matrix of type '<class 'numpy.float64'>'  
with 100226 stored elements in COOrdinate format>
```

```
data.toarray() #This is how we can see the data visually
```

```
array([[0.83333333, 0.83333333, 0.83333333, ..., 0.          , 0.          ,  
       0.          ],  
     [0.          , 0.          , 0.          , ..., 0.          , 0.          ,  
      0.          ],  
     [0.          , 0.          , 0.          , ..., 0.          , 0.          ,  
      0.          ],  
     ...,  
     [0.5          , 0.5          , 0.          , ..., 0.          , 0.          ,  
      0.          ],  
     [0.66666667, 0.          , 0.          , ..., 0.          , 0.          ,  
      0.          ],  
     [1.          , 0.          , 1.          , ..., 0.66666667, 0.66666667,  
      0.66666667]])
```

```
n = data.get_shape()[0]  
m = data.get_shape()[1]  
avg = data.sum()/data.count_nonzero()  
print(f'Data has {n} rows and {m} columns')  
print(f'Average rating is {avg:.4F}')
```

```
Data has 610 rows and 9724 columns  
Average rating is 0.7249
```

```
#Similarly, we prepare the validation data matrix  
tmp = np.genfromtxt('ratings_validation.csv', delimiter=',', dtype='int32')  
val = np.array(tmp[:,2]+1, dtype='float')/6  
validation = coo_matrix( (val, (tmp[:,0]-1, tmp[:,1]-1)) )
```

```
validation
```

```
<610x9341 sparse matrix of type '<class 'numpy.float64'>'  
with 610 stored elements in COOrdinate format>
```

```
#SGD is provided in the additional file util.py which has been uploaded  
# Compute just factors  
res = sgd(data, 10, factor_learning_rate=0.001, epochs=100, verbose=True ) #factor_learning_r
```

```
Epoch 41: error 3098.03      factor_learning rate = 0.001  
Epoch 42: error 3053.61      factor_learning rate = 0.001  
Epoch 43: error 3012.08      factor_learning rate = 0.001  
Epoch 44: error 2970.45      factor_learning rate = 0.001  
Epoch 45: error 2931.56      factor_learning rate = 0.001  
Epoch 46: error 2890.29      factor_learning rate = 0.001  
Epoch 47: error 2856.78      factor_learning rate = 0.001  
Epoch 48: error 2821.09      factor_learning rate = 0.001  
Epoch 49: error 2784.41      factor_learning rate = 0.001  
Epoch 50: error 2753.31      factor_learning rate = 0.001  
Epoch 51: error 2719.67      factor_learning rate = 0.001  
Epoch 52: error 2688.62      factor_learning rate = 0.001  
Epoch 53: error 2658.55      factor_learning rate = 0.001  
Epoch 54: error 2629.06      factor_learning rate = 0.001  
Epoch 55: error 2601.25      factor_learning rate = 0.001  
Epoch 56: error 2571.48      factor_learning rate = 0.001  
Epoch 57: error 2545.28      factor_learning rate = 0.001  
Epoch 58: error 2518.87      factor_learning rate = 0.001  
Epoch 59: error 2495.94      factor_learning rate = 0.001  
Epoch 60: error 2471.06      factor_learning rate = 0.001  
Epoch 61: error 2445.79      factor_learning rate = 0.001  
Epoch 62: error 2421.83      factor_learning rate = 0.001  
Epoch 63: error 2401.31      factor_learning rate = 0.001  
Epoch 64: error 2377.63      factor_learning rate = 0.001  
Epoch 65: error 2355.29      factor_learning rate = 0.001  
Epoch 66: error 2336.52      factor_learning rate = 0.001  
Epoch 67: error 2314.03      factor_learning rate = 0.001  
Epoch 68: error 2296.60      factor_learning rate = 0.001  
Epoch 69: error 2276.61      factor_learning rate = 0.001  
Epoch 70: error 2257.62      factor_learning rate = 0.001  
Epoch 71: error 2239.59      factor_learning rate = 0.001  
Epoch 72: error 2223.42      factor_learning rate = 0.001  
Epoch 73: error 2204.06      factor_learning rate = 0.001  
Epoch 74: error 2188.74      factor_learning rate = 0.001  
Epoch 75: error 2172.01      factor_learning rate = 0.001  
Epoch 76: error 2157.77      factor_learning rate = 0.001  
Epoch 77: error 2139.62      factor_learning rate = 0.001  
Epoch 78: error 2126.47      factor_learning rate = 0.001  
Epoch 79: error 2110.57      factor_learning rate = 0.001  
Epoch 80: error 2094.65      factor_learning rate = 0.001  
Epoch 81: error 2082.63      factor_learning rate = 0.001  
Epoch 82: error 2067.80      factor_learning rate = 0.001  
Epoch 83: error 2053.95      factor_learning rate = 0.001  
Epoch 84: error 2041.58      factor_learning rate = 0.001  
Epoch 85: error 2028.66      factor_learning rate = 0.001  
Epoch 86: error 2015.64      factor_learning rate = 0.001  
Epoch 87: error 2003.03      factor_learning rate = 0.001  
Epoch 88: error 1991.94      factor_learning rate = 0.001  
Epoch 89: error 1979.44      factor_learning rate = 0.001
```

```

Epoch  90: error 1969.31      factor_learning_rate = 0.001
Epoch  91: error 1958.02      factor_learning_rate = 0.001
Epoch  92: error 1948.35      factor_learning_rate = 0.001
Epoch  93: error 1936.20      factor_learning_rate = 0.001
Epoch  94: error 1927.27      factor_learning_rate = 0.001
Epoch  95: error 1915.00      factor_learning_rate = 0.001
Epoch  96: error 1904.43      factor_learning_rate = 0.001
Epoch  97: error 1895.47      factor_learning_rate = 0.001
Epoch  98: error 1885.35      factor_learning_rate = 0.001
Epoch  99: error 1877.05      factor_learning_rate = 0.001

```

#We can see by the verbose that the error keeps on decreasing

```

val_err = squared_error(validation, res.L, res.R)
print('Only factors, fixed rate, no regularization')
print(f'Final training error = {res.err[-1]:5.3f}')
    f'\tRelative error = {res.err[-1]/squared_error(data)}'
print(f'Validation error = {val_err:5.3f}'
    f'\tRelative error = {val_err/squared_error(validation)}')

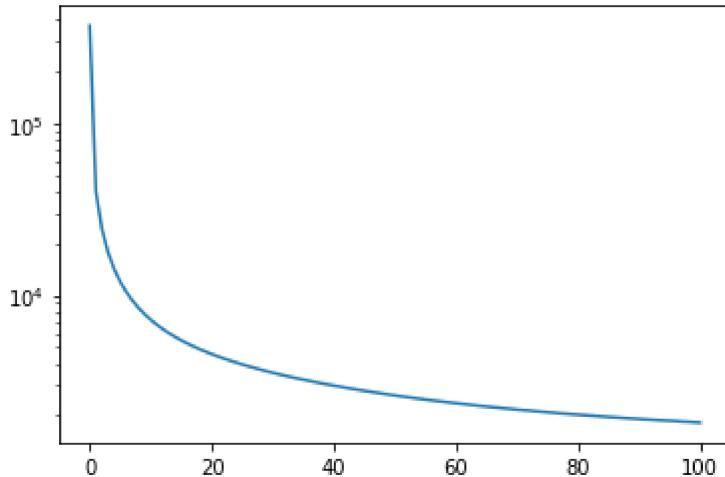
Only factors, fixed rate, no regularization
Final training error = 1817.814 Relative error = 0.032481741082324427
Validation error = 229.815  Relative error = 0.6027504825942747

```

```

# Plot errors over epochs
plt.plot(res.err)
plt.gca().set_yscale('log') # log scale makes these plots easier to read
plt.show()
plt.close()

```



#This confirms that the error keeps on decreasing as we saw in verbose

#However, the error has not converged at 100 epochs given that the curve values are still decreasing.
#we can agree to a convergence value when the difference is very small.

```
#The algorithm does not do early stopping if the change between two
#consecutive epochs is small enough. In stochastic gradient descent
#since the gradient vector is chosen randomly and so it is very much possible
#that the algorithm might not reach at exact global minima. Thus, it would be
#very much sensible to stop at a a minimum difference in consecutive epoch value
```

```
# Read movies
movies = []
with open('movies.txt') as f:
    for row in f:
        movies.append(row.strip())

# Print the top-10 and bottom-10 movies
print_factors(res.R, movies)

2 Caddyshack II (1988) -0.240
3 Curly Sue (1991) -0.229
4 Always Watching: A Marble Hornets Story (2015) -0.190
5 Assassin's Creed (2016) -0.183
6 Candyman 3: Day of the Dead (1999) -0.181
7 Hard Ticket to Hawaii (1987) -0.176
8 Detroit Rock City (1999) -0.174
9 What the #$*! Do We Know!? (a.k.a. What the Bleep Do We Know!?) (2004) -0.16
10 Vie en Rose, La (Môme, La) (2007) -0.169
```

Factor 8 top-10

```
1 Dam Busters, The (1955) 0.374
2 Cincinnati Kid, The (1965) 0.364
3 Infinity (1996) 0.362
4 Used People (1992) 0.361
5 Thin Man Goes Home, The (1945) 0.361
6 Sea of Love (1989) 0.360
7 Angst (1983) 0.360
8 In Love and War (1996) 0.360
9 Losin' It (1983) 0.359
10 Jazz Singer, The (1927) 0.358
```

Factor 8 bottom-10

```
1 Baby Geniuses (1999) -0.212
2 Angel Heart (1987) -0.208
3 Unknown (2006) -0.202
4 They Shoot Horses, Don't They? (1969) -0.200
5 Wagons East (1994) -0.190

6 Edge of Heaven, The (Auf der anderen Seite) (2007) -0.186
7 Asphyx, The (1973) -0.180
8 Babylon 5 -0.179
9 Whale Rider (2002) -0.176
10 Devil and Daniel Johnston, The (2005) -0.173
```

Factor 9 top-10

```
1 Comic-Con Episode IV: A Fan's Hope (2011) 0.382
2 Returner (Ritaanaa) (2002) 0.381
3 Revolution (1985) 0.375
4 Sleepwalkers (1992) 0.374
5 The Flash 2 - Revenge of the Trickster (1991) 0.373
```

```

6    Afro Samurai (2007)      0.371
7    Mutant Aliens (2001)     0.367
8    V. I. Warshawski (1991)  0.366
9    Lionheart (1990)        0.365
10   Danny Deckchair (2003)   0.365

```

Factor 9 bottom-10

```

1    Amityville: A New Generation (1993)      -0.254
2    Doug's 1st Movie (1999)                   -0.246
3    Loverboy (1989)                         -0.227
4    The Star Wars Holiday Special (1978)      -0.221
5    Pollyanna (1960)                        -0.211
6    Ringu 0: Bâsudei (2000)                  -0.209
7    Turbo: A Power Rangers Movie (1997)      -0.202
8    La vérité si je mens ! (1997)           -0.200
9    Strangeland (1998)                      -0.195
10   Newton Boys, The (1998)                 -0.194

```

```
#The pattern is that the top movies have score above 0.3 while bottom ones have scores
# in negative
```

```
# Compute just factors, but use bold driver and slower learning
res = sgd(data, 10, bold_driver=2, factor_learning_rate=0.0001)
```

```

Epoch 41: error 1411.00          factor_learning_rate = 0.0001
Epoch 42: error 1414.60          factor_learning_rate = 0.0002
Epoch 43: error 1412.80          factor_learning_rate = 0.0004
Epoch 44: error 1412.27          factor_learning_rate = 0.0008
Epoch 45: error 1414.10          factor_learning_rate = 0.0001
Epoch 46: error 1411.39          factor_learning_rate = 0.0002
Epoch 47: error 1409.80          factor_learning_rate = 0.0004
Epoch 48: error 1409.35          factor_learning_rate = 0.0008
Epoch 49: error 1411.00          factor_learning_rate = 0.0001
Epoch 50: error 1408.58          factor_learning_rate = 0.0002
Epoch 51: error 1407.08          factor_learning_rate = 0.0004
Epoch 52: error 1406.59          factor_learning_rate = 0.0008
Epoch 53: error 1408.69          factor_learning_rate = 0.0001
Epoch 54: error 1406.12          factor_learning_rate = 0.0002
Epoch 55: error 1404.52          factor_learning_rate = 0.0004
Epoch 56: error 1404.16          factor_learning_rate = 0.0008
Epoch 57: error 1406.26          factor_learning_rate = 0.0001
Epoch 58: error 1403.55          factor_learning_rate = 0.0002
Epoch 59: error 1402.00          factor_learning_rate = 0.0004
Epoch 60: error 1401.61          factor_learning_rate = 0.0008
Epoch 61: error 1403.09          factor_learning_rate = 0.0001
Epoch 62: error 1400.81          factor_learning_rate = 0.0002
Epoch 63: error 1399.53          factor_learning_rate = 0.0004
Epoch 64: error 1399.33          factor_learning_rate = 0.0008
Epoch 65: error 1402.02          factor_learning_rate = 0.0001
Epoch 66: error 1398.83          factor_learning_rate = 0.0002
Epoch 67: error 1397.27          factor_learning_rate = 0.0004
Epoch 68: error 1397.07          factor_learning_rate = 0.0008
Epoch 69: error 1398.37          factor_learning_rate = 0.0001
Epoch 70: error 1396.30          factor_learning_rate = 0.0002

```

```

Epoch 71: error 1394.99      factor_learning rate = 0.0004
Epoch 72: error 1394.69      factor_learning rate = 0.0008
Epoch 73: error 1396.56      factor_learning rate = 0.0001
Epoch 74: error 1394.22      factor_learning rate = 0.0002
Epoch 75: error 1392.76      factor_learning rate = 0.0004
Epoch 76: error 1392.53      factor_learning rate = 0.0008
Epoch 77: error 1394.88      factor_learning rate = 0.0001
Epoch 78: error 1392.23      factor_learning rate = 0.0002
Epoch 79: error 1390.66      factor_learning rate = 0.0004
Epoch 80: error 1390.36      factor_learning rate = 0.0008
Epoch 81: error 1392.31      factor_learning rate = 0.0001
Epoch 82: error 1389.88      factor_learning rate = 0.0002
Epoch 83: error 1388.38      factor_learning rate = 0.0004
Epoch 84: error 1388.20      factor_learning rate = 0.0008
Epoch 85: error 1389.61      factor_learning rate = 0.0001

Epoch 86: error 1387.55      factor_learning rate = 0.0002
Epoch 87: error 1386.33      factor_learning rate = 0.0004
Epoch 88: error 1385.97      factor_learning rate = 0.0008
Epoch 89: error 1388.05      factor_learning rate = 0.0001
Epoch 90: error 1385.76      factor_learning rate = 0.0002
Epoch 91: error 1384.35      factor_learning rate = 0.0004
Epoch 92: error 1383.98      factor_learning rate = 0.0008
Epoch 93: error 1385.41      factor_learning rate = 0.0001
Epoch 94: error 1383.56      factor_learning rate = 0.0002
Epoch 95: error 1382.23      factor_learning rate = 0.0004
Epoch 96: error 1381.97      factor_learning rate = 0.0008
Epoch 97: error 1384.69      factor_learning rate = 0.0001
Epoch 98: error 1381.82      factor_learning rate = 0.0002
Epoch 99: error 1380.31      factor_learning rate = 0.0004

```

```

val_err = squared_error(validation, res.L, res.R)
print('Only factors, fixed rate, no regularization')
print(f'Final training error = {res.err[-1]:5.3f}')
    f'\tRelative error = {res.err[-1]/squared_error(data)}'
print(f'Validation error = {val_err:5.3f}'
    f'\tRelative error = {val_err/squared_error(validation)}')

```

```

Only factors, fixed rate, no regularization
Final training error = 1380.074 Relative error = 0.024659964429735608
    Validation error = 221.044  Relative error = 0.5797447581574955

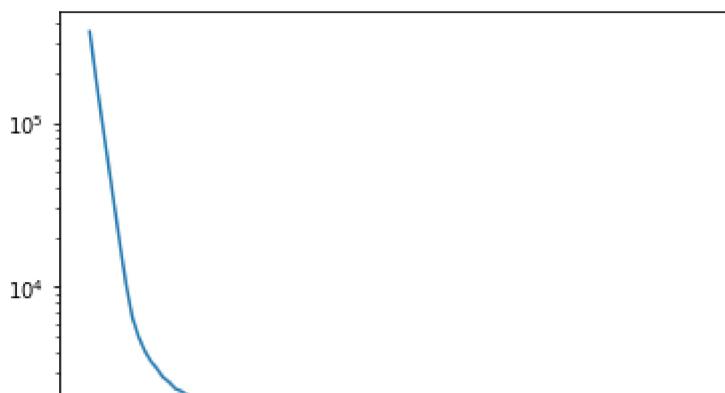
```

#So we see that the validation error has reduced this time with bold driver heuristic

```

# Plot errors over epochs
plt.plot(res.err)
plt.gca().set_yscale('log') # log scale makes these plots easier to read
plt.show()
plt.close()

```



#The plot above clearly shows that the convergence to a more straight line takes place
#much faster this time

```
# Compute just factors, but use bold driver and slower learning
```

```
res = sgd(data, 10, bold_driver=3, factor_learning_rate=0.001)
```

```
Epoch 41: error 1558.55      factor_learning_rate = 0.003
Epoch 42: error 1385.69      factor_learning_rate = 0.001
Epoch 43: error 1354.12      factor_learning_rate = 0.003
Epoch 44: error 1382.99      factor_learning_rate = 0.001
Epoch 45: error 1349.12      factor_learning_rate = 0.003
Epoch 46: error 1377.30      factor_learning_rate = 0.001
Epoch 47: error 1346.17      factor_learning_rate = 0.003
Epoch 48: error 1378.10      factor_learning_rate = 0.001
Epoch 49: error 1341.68      factor_learning_rate = 0.003
Epoch 50: error 1371.94      factor_learning_rate = 0.001
Epoch 51: error 1339.61      factor_learning_rate = 0.003
Epoch 52: error 1370.26      factor_learning_rate = 0.001
Epoch 53: error 1333.97      factor_learning_rate = 0.003
Epoch 54: error 1359.49      factor_learning_rate = 0.001
Epoch 55: error 1330.27      factor_learning_rate = 0.003
Epoch 56: error 1362.01      factor_learning_rate = 0.001
Epoch 57: error 1327.28      factor_learning_rate = 0.003
Epoch 58: error 1358.79      factor_learning_rate = 0.001
Epoch 59: error 1324.17      factor_learning_rate = 0.003
Epoch 60: error 1349.38      factor_learning_rate = 0.001
Epoch 61: error 1320.25      factor_learning_rate = 0.003
Epoch 62: error 1348.73      factor_learning_rate = 0.001
Epoch 63: error 1317.52      factor_learning_rate = 0.003
Epoch 64: error 1345.52      factor_learning_rate = 0.001
Epoch 65: error 1314.02      factor_learning_rate = 0.003
Epoch 66: error 1344.71      factor_learning_rate = 0.001
Epoch 67: error 1311.98      factor_learning_rate = 0.003
Epoch 68: error 1339.31      factor_learning_rate = 0.001
Epoch 69: error 1307.83      factor_learning_rate = 0.003
Epoch 70: error 1335.77      factor_learning_rate = 0.001
Epoch 71: error 1305.64      factor_learning_rate = 0.003
Epoch 72: error 1336.85      factor_learning_rate = 0.001
Epoch 73: error 1302.57      factor_learning_rate = 0.003
Epoch 74: error 1329.09      factor_learning_rate = 0.001
Epoch 75: error 1299.83      factor_learning_rate = 0.003
Epoch 76: error 1331.98      factor_learning_rate = 0.001
Epoch 77: error 1296.26      factor_learning_rate = 0.003
Epoch 78: error 1327.48      factor_learning_rate = 0.001
Epoch 79: error 1305.47      factor_learning_rate = 0.001
```

```

Epoch  79: error 1295.45      factor_learning_rate = 0.003
Epoch  80: error 1320.53      factor_learning_rate = 0.001
Epoch  81: error 1291.63      factor_learning_rate = 0.003
Epoch  82: error 1323.90      factor_learning_rate = 0.001
Epoch  83: error 1289.14      factor_learning_rate = 0.003
Epoch  84: error 1320.91      factor_learning_rate = 0.001
Epoch  85: error 1287.71      factor_learning_rate = 0.003
Epoch  86: error 1315.20      factor_learning_rate = 0.001
Epoch  87: error 1284.06      factor_learning_rate = 0.003
Epoch  88: error 1314.30      factor_learning_rate = 0.001
Epoch  89: error 1282.15      factor_learning_rate = 0.003
Epoch  90: error 1310.41      factor_learning_rate = 0.001
Epoch  91: error 1279.72      factor_learning_rate = 0.003
Epoch  92: error 1310.53      factor_learning_rate = 0.001
Epoch  93: error 1278.29      factor_learning_rate = 0.003

Epoch  94: error 1307.28      factor_learning_rate = 0.001
Epoch  95: error 1275.29      factor_learning_rate = 0.003
Epoch  96: error 1303.02      factor_learning_rate = 0.001
Epoch  97: error 1273.62      factor_learning_rate = 0.003
Epoch  98: error 1301.70      factor_learning_rate = 0.001
Epoch  99: error 1271.54      factor_learning_rate = 0.003

```

```

val_err = squared_error(validation, res.L, res.R)
print('Only factors, fixed rate, no regularization')
print(f'Final training error = {res.err[-1]:5.3f}')
    f'\tRelative error = {res.err[-1]/squared_error(data)}'
print(f'    Validation error = {val_err:5.3f}'
    f'\tRelative error = {val_err/squared_error(validation)})')

Only factors, fixed rate, no regularization
Final training error = 1297.298 Relative error = 0.023180870199845993
    Validation error = 191.577  Relative error = 0.5024611010341171

```

#So now I manage to reduce the error further

```

# Plot errors over epochs
plt.plot(res.err)
plt.gca().set_yscale('log') # log scale makes these plots easier to read
plt.show()
plt.close()

```

```
# Compute just factors, but use bold driver and slower learning
res = sgd(data, 10, bold_driver=3, factor_learning_rate=0.01)
```

```

Epoch  91: error 1358.28      factor_learning_rate = 0.01
Epoch  92: error 1349.60      factor_learning_rate = 0.01
Epoch  93: error 1361.25      factor_learning_rate = 0.01
Epoch  94: error 1356.77      factor_learning_rate = 0.01
Epoch  95: error 1369.15      factor_learning_rate = 0.01
Epoch  96: error 1370.30      factor_learning_rate = 0.01
Epoch  97: error 1331.42      factor_learning_rate = 0.01
Epoch  98: error 1347.75      factor_learning_rate = 0.01
Epoch  99: error 1352.07      factor_learning_rate = 0.01

```

```

val_err = squared_error(validation, res.L, res.R)
print('Only factors, fixed rate, no regularization')
print(f'Final training error = {res.err[-1]:5.3f}')
    f'\tRelative error = {res.err[-1]/squared_error(data)}'
print(f'Validation error = {val_err:5.3f}')
    f'\tRelative error = {val_err/squared_error(validation)}')

Only factors, fixed rate, no regularization
Final training error = 1365.280 Relative error = 0.024395606035237705
Validation error = 188.997  Relative error = 0.49569468928693805

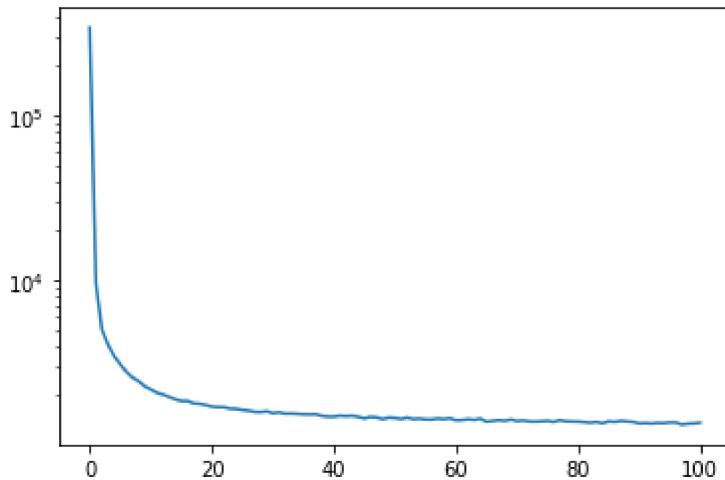
```

#The validation error drops further :)

```

# Plot errors over epochs
plt.plot(res.err)
plt.gca().set_yscale('log') # log scale makes these plots easier to read
plt.show()
plt.close()

```



```

# Compute factors and bias
res = sgd(data, 10, bold_driver=2, factor_learning_rate=0.0001, bias_learning_rate=0.0001, ve
Epoch  41: error 1315.61      factor_learning_rate = 0.0008
Epoch  42: error 1317.74      factor_learning_rate = 0.0001
Epoch  43: error 1315.01      factor_learning_rate = 0.0002
Epoch  44: error 1313.44      factor_learning_rate = 0.0004
Epoch  45: error 1313.07      factor_learning_rate = 0.0008
Epoch  46: error 1315.36      factor_learning_rate = 0.0001

```

```
val_err = squared_error(validation, res.L, res.R, res.global_bias, res.row_bias, res.col_bias)
print('Bias only, no regularization')
```

```
print(f'Final training error = {res.err[-1]:5.3f}')
    f'\tRelative error = {res.err[-1]/squared_error(data)}')
print(f'Validation error = {val_err:5.3f}')
    f'\tRelative error = {val_err/squared_error(validation)}')

Bias only, no regularization
Final training error = 1285.582 Relative error = 0.02297151945956205
Validation error = 222.025  Relative error = 0.5823179086533342

# Print the movies based on their bias top 10 and bottom 10
print_factors(res.col_bias, movies)

Factor 0 top-10
1 Pulp Fiction (1994) 0.299
2 Cell 211 (Celda 211) (2009) 0.279
3 Reservoir Dogs (1992) 0.264
4 Hearts of Darkness: A Filmmakers Apocalypse (1991) 0.259
5 Watchmen: Tales of the Black Freighter (2009) 0.257
6 Syrup (2013) 0.255
7 Stalker (1979) 0.248
8 Léon: The Professional (a.k.a. The Professional) (Léon) (1994) 0.241
9 And the Band Played On (1993) 0.236
10 Fist Fight (2017) 0.235

Factor 0 bottom-10
1 No Small Affair (1984) -0.455
2 Infernal Affairs 2 (Mou gaan dou II) (2003) -0.388
3 Stefan Zweig: Farewell to Europe (2016) -0.362
4 Urban Legends: Final Cut (2000) -0.353
5 Grind (2003) -0.351
6 Aelita: The Queen of Mars (Aelita) (1924) -0.347
7 My Bloody Valentine (1981) -0.344
8 Boat Trip (2003) -0.344
9 The Crew (2016) -0.343
10 3 Days to Kill (2014) -0.343

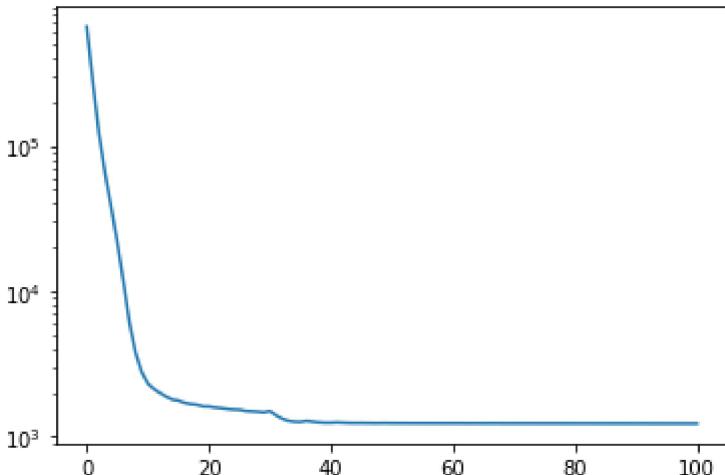
# Plot errors over epochs
plt.plot(res.err)
plt.gca().set_yscale('log') # log scale makes these plots easier to read
plt.show()
plt.close()
#The results aren't better because the validation error is 222 while I was already
#able to achieve a validation error of 188 earlier.
```



```
#Now doing SGD with regularization L2-Regularization
res = sgd(data, 10, bold_driver=2, factor_learning_rate=0.0001, bias_learning_rate=0.0001, f=
           |   |
val_err = squared_error(validation, res.L, res.R, res.global_bias, res.row_bias, res.col_bias)
print('Factors and bias, regularized')
print(f'Final training error = {res.err[-1]:5.3f}')
    f'\tRelative error = {res.err[-1]/squared_error(data)}'
print(f'Validation error = {val_err:5.3f}'
    f'\tRelative error = {val_err/squared_error(validation)}')

Factors and bias, regularized
Final training error = 1231.810 Relative error = 0.022010698393192092
    Validation error = 203.652  Relative error = 0.5341297631399506

# Plot errors over epochs
plt.plot(res.err)
plt.gca().set_yscale('log') # log scale makes these plots easier to read
plt.show()
plt.close()
#The results aren't better because the validation error is 222 while I was already
#able to achieve a validation error of 188 earlier.
```



```
regularization L2-Regularization, bold driver of 3, factor learning rate of 0.01
old_driver=3, factor_learning_rate=0.01, bias_learning_rate=0.0001, factor_regu=0.001, bias_r

Epoch  0: error 653219.96      factor_learning_rate = 0.01
Epoch  1: error 8972.64        factor_learning_rate = 0.01
Epoch  2: error 4457.27        factor_learning_rate = 0.01
Epoch  3: error 3431.22        factor_learning_rate = 0.01
```

```
Epoch  4: error 2870.53      factor_learning_rate = 0.01
Epoch  5: error 2546.66      factor_learning_rate = 0.01
Epoch  6: error 2306.88      factor_learning_rate = 0.01
Epoch  7: error 2141.87      factor_learning_rate = 0.01
Epoch  8: error 2019.36      factor_learning_rate = 0.01
Epoch  9: error 1932.13      factor_learning_rate = 0.01
Epoch 10: error 1866.20      factor_learning_rate = 0.01
Epoch 11: error 1807.71      factor_learning_rate = 0.01
Epoch 12: error 1771.94      factor_learning_rate = 0.01
Epoch 13: error 1712.39      factor_learning_rate = 0.01
Epoch 14: error 1667.06      factor_learning_rate = 0.01
Epoch 15: error 1640.71      factor_learning_rate = 0.01
Epoch 16: error 1613.60      factor_learning_rate = 0.01
Epoch 17: error 1586.59      factor_learning_rate = 0.01
Epoch 18: error 1564.48      factor_learning_rate = 0.01
Epoch 19: error 1537.26      factor_learning_rate = 0.01
Epoch 20: error 1533.13      factor_learning_rate = 0.01
Epoch 21: error 1528.86      factor_learning_rate = 0.01
Epoch 22: error 1502.74      factor_learning_rate = 0.01
Epoch 23: error 1490.45      factor_learning_rate = 0.01
Epoch 24: error 1469.40      factor_learning_rate = 0.01
Epoch 25: error 1456.77      factor_learning_rate = 0.01
Epoch 26: error 1455.65      factor_learning_rate = 0.01
Epoch 27: error 1426.81      factor_learning_rate = 0.01
Epoch 28: error 1412.73      factor_learning_rate = 0.01
Epoch 29: error 1409.39      factor_learning_rate = 0.01
Epoch 30: error 1404.42      factor_learning_rate = 0.01
Epoch 31: error 1399.30      factor_learning_rate = 0.01
Epoch 32: error 1383.29      factor_learning_rate = 0.01
Epoch 33: error 1384.88      factor_learning_rate = 0.01
Epoch 34: error 1355.22      factor_learning_rate = 0.01
Epoch 35: error 1351.55      factor_learning_rate = 0.01
Epoch 36: error 1343.42      factor_learning_rate = 0.01
Epoch 37: error 1340.98      factor_learning_rate = 0.01
Epoch 38: error 1352.54      factor_learning_rate = 0.01
Epoch 39: error 1332.35      factor_learning_rate = 0.01
Epoch 40: error 1322.87      factor_learning_rate = 0.01
Epoch 41: error 1307.07      factor_learning_rate = 0.01
Epoch 42: error 1308.47      factor_learning_rate = 0.01
Epoch 43: error 1315.10      factor_learning_rate = 0.01
Epoch 44: error 1294.51      factor_learning_rate = 0.01
Epoch 45: error 1310.83      factor_learning_rate = 0.01
Epoch 46: error 1292.40      factor_learning_rate = 0.01
Epoch 47: error 1283.86      factor_learning_rate = 0.01
Epoch 48: error 1290.73      factor_learning_rate = 0.01
Epoch 49: error 1280.16      factor_learning_rate = 0.01
Epoch 50: error 1278.88      factor_learning_rate = 0.01
Epoch 51: error 1278.71      factor_learning_rate = 0.01
Epoch 52: error 1280.26      factor_learning_rate = 0.01
Epoch 53: error 1262.37      factor_learning_rate = 0.01
Epoch 54: error 1268.39      factor_learning_rate = 0.01
Epoch 55: error 1263.37      factor_learning_rate = 0.01
Epoch 56: error 1266.02      factor_learning_rate = 0.01
Epoch 57: error 1259.22      factor_learning_rate = 0.01
Epoch 58: error 1254.53      factor_learning_rate = 0.01
```

```
val_err = squared_error(validation, res.L, res.R, res.global_bias, res.row_bias, res.col_bias)
```

```

print( factors_and_bias, regularized )
print(f'Final training error = {res.err[-1]:5.3f}'
      f'\tRelative error = {res.err[-1]/squared_error(data)}')
print(f'Validation error = {val_err:5.3f}'
      f'\tRelative error = {val_err/squared_error(validation)})')

```

Factors and bias, regularized

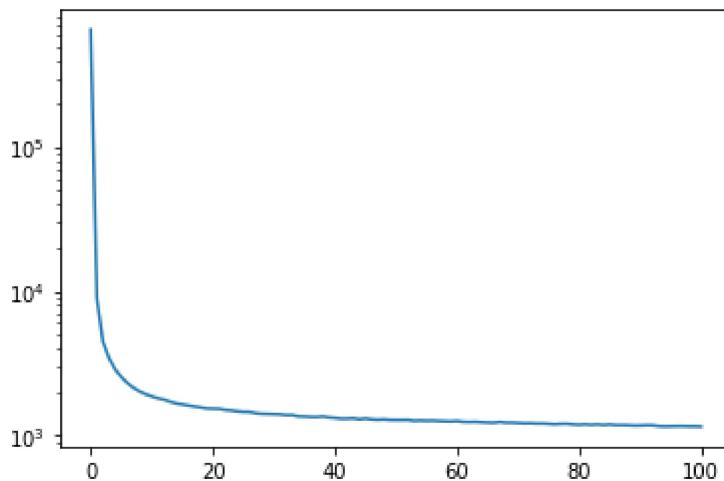
Final training error = 1151.517 Relative error = 0.020575962304935237
 Validation error = 267.460 Relative error = 0.7014830460905507

Plot errors over epochs

```

plt.plot(res.err)
plt.gca().set_yscale('log') # log scale makes these plots easier to read
plt.show()
plt.close()

```



#Regularization comes with a cost of increasing the error rate. The validation error has incr
#But perhaps that is what is needed.

```
fullArray = data.toarray()
```

```

from numpy.linalg import svd
## Compute SVD of data
U, S, V = svd(fullArray, full_matrices=False, compute_uv=True)

```

#Squared Error

```
se = S @ V
```

```
se
```

```
array([ 7.90568756, -8.37391345, -0.05052415, ... , -0.73988401,
       -0.73988401, -0.73988401])
```

```
#To get truncated SVD
from sklearn.utils.extmath import randomized_svd

Utruncated, Sigma, VTtruncated = randomized_svd(fullArray,n_components=10,n_iter=5,random_state=None)

seTruncated = Sigma @ VTtruncated

seTruncated

array([0.06608784, 0.64533515, 4.51407857, ..., 0.26432449, 0.26432449,
       0.26432449])

#Now doing the same for validation data

validationFull = validation.toarray()

validationFull

array([[0., 0., 0., ..., 0., 0., 0.],
       [0., 0., 0., ..., 0., 0., 0.],
       [0., 0., 0., ..., 0., 0., 0.],
       ...,
       [0., 0., 0., ..., 0., 0., 0.],
       [0., 0., 0., ..., 0., 0., 0.],
       [0., 0., 0., ..., 0., 0., 0.]])
```



```
UtruncatedValid, SigmaValid, VTtruncatedValid = randomized_svd(validationFull,n_components=10,random_state=None)

seValidationTruncated = SigmaValid @ VTtruncatedValid

seValidationTruncated

array([-1.66403198e-16, -2.38323818e-16, 3.52371195e-16, ...,
       0.00000000e+00, 0.00000000e+00, -9.59328450e-07])
```



```
len(seValidationTruncated)
```



```
9341
```



```
len(seTruncated)
```



```
9724
```

```
#Since the lenghts are differenct, they cannot be 1 to 1 compared.
```

```
#Task 2#
```

```
#Task 2: Initial solutions and hyperparameter optimization
```

```
#We continue using the same data and code as above. We start by looking at the effect of the
#solutions. By default, the code uses uniform distribution from unit interval for factor matr
#normalized, and normal distribution with 0.1 standard deviation for biases. Try following c
#solutions and report the results. Explain what happened and why you think it happened (e.g.
#better because something or the code crashed because of something else).
```

```
#a) Factor matrices and bias vectors are all-1s matrices and vectors. (No normalization.)
```

```
#b) Factor matrices and bias vectors are all-0s matrices and vectors. (No normalization.)
```

```
#c) Factor matrices are drawn from normal distribution with 0.1 standard deviation. (No norma
```

```
#d) Factor matrices are random integer matrices with values from 0 to 6, normalized.
```

```
#Next we try to optimize the hyperparameters. The algorithm has the following hyperparameters
```

```
#a) rank
```

```
#b) number of epochs
```

```
#c) learning rates (for simplicity, let's assume they're same for factors and biases)
```

```
#d) bold driver multiplier
```

```
#e) regularization coefficients (again, same for factors and biases).
```

```
#
```

```
#The standard way to optimize these is to select some values for each hyperparameter and try
#combinations of the potential values. This is called a grid search. To test how good each cc
#one should also run cross validation; a 5-fold cross validation, for example, would split th
#into five random subsets, find the decomposition based on four of them, and test the result
#repeat four more times so that each subset would get left out once. The final quality value
#average error over the five runs. Assume we want to try 5 different values for each hyperpar
#5-fold cross validation. Time the algorithm running once over the data set and estimate how
#grid search would take. You can assume 5-fold cross validation takes five times the time one
#that run time grows linearly with number of epochs and rank. Report what you calculated.
```

```
#If the expected run time exceeds the time you have before the DL, design an experiment where
#at least the following (you can change the values if you want, but explain why in your repor
#a) rank = 10
```

```
#b) number of epochs = 100 (or smaller if you think you converge)
```

```
#c) learning rates in {10-5
```

```
, 10-4
```

```
, 10-3}
```

```
#d) bold driver multiplier = 2 (or some other coefficient)
```

```
#e) regularization coefficients in {10-6
```

```
, 10-5
```

```
, 10-4}.
```

```
#Notice that this still leaves you with nine combinations that you must be able to run. Be pr
#your computer running overnight!
```

```
#After you've found the optimum combination of hyperparameters, run SGD with full training da
#those parameters and compute the validation error. Is the error better than the best you've
#Argue!
```

```
#a) Factor matrices and bias vectors are all-1s matrices and vectors. (No normalization.)
```

```
n = data.get_shape()[0]
```

```
m = data.get_shape()[1]
```

```
avg = data.sum()/data.count_nonzero()
```

```
# Initial factors are given as a tuple (L, R); R is transposed
init_LR = (np.ones(shape=(n,10)), np.ones(shape=(m,10)))
# Initial biases are a tuple (global_bias, row_bias, col_bias)
# Give the true global bias.
init_biases = (avg, np.ones(shape=n), np.ones(shape=m))

# You can change the learning rates and bold driver factor if you want
res = sgd(data, 10, bold_driver=2, factor_learning_rate=0.0001, bias_learning_rate=0.0001, ve
```

Epoch 0:	error 14435843.09	factor_learning rate = 0.0001
Epoch 1:	error 3687950.83	factor_learning rate = 0.0002
Epoch 2:	error 803754.06	factor_learning rate = 0.0004
Epoch 3:	error 327210.90	factor_learning rate = 0.0008
Epoch 4:	error 143830.17	factor_learning rate = 0.0016
Epoch 5:	error 63452.22	factor_learning rate = 0.0032
Epoch 6:	error 27742.48	factor_learning rate = 0.0064
Epoch 7:	error 12648.93	factor_learning rate = 0.01
Epoch 8:	error 7855.12	factor_learning rate = 0.01
Epoch 9:	error 4951.88	factor_learning rate = 0.01
Epoch 10:	error 3806.94	factor_learning rate = 0.01
Epoch 11:	error 3460.73	factor_learning rate = 0.01
Epoch 12:	error 3345.87	factor_learning rate = 0.01
Epoch 13:	error 3127.46	factor_learning rate = 0.01
Epoch 14:	error 3115.84	factor_learning rate = 0.01
Epoch 15:	error 2973.42	factor_learning rate = 0.01
Epoch 16:	error 2926.97	factor_learning rate = 0.01
Epoch 17:	error 2875.61	factor_learning rate = 0.01
Epoch 18:	error 2770.24	factor_learning rate = 0.01
Epoch 19:	error 2828.46	factor_learning rate = 0.0001
Epoch 20:	error 2355.66	factor_learning rate = 0.0002
Epoch 21:	error 2122.54	factor_learning rate = 0.0004
Epoch 22:	error 2001.82	factor_learning rate = 0.0008
Epoch 23:	error 1947.30	factor_learning rate = 0.0016
Epoch 24:	error 1947.04	factor_learning rate = 0.0032
Epoch 25:	error 2009.12	factor_learning rate = 0.0001
Epoch 26:	error 1940.80	factor_learning rate = 0.0002
Epoch 27:	error 1910.30	factor_learning rate = 0.0004
Epoch 28:	error 1899.52	factor_learning rate = 0.0008
Epoch 29:	error 1902.07	factor_learning rate = 0.0001
Epoch 30:	error 1891.75	factor_learning rate = 0.0002
Epoch 31:	error 1889.24	factor_learning rate = 0.0004
Epoch 32:	error 1890.69	factor_learning rate = 0.0001
Epoch 33:	error 1887.37	factor_learning rate = 0.0002
Epoch 34:	error 1886.69	factor_learning rate = 0.0004
Epoch 35:	error 1888.54	factor_learning rate = 0.0001
Epoch 36:	error 1885.73	factor_learning rate = 0.0002
Epoch 37:	error 1885.34	factor_learning rate = 0.0004
Epoch 38:	error 1887.79	factor_learning rate = 0.0001
Epoch 39:	error 1884.48	factor_learning rate = 0.0002
Epoch 40:	error 1884.36	factor_learning rate = 0.0004
Epoch 41:	error 1887.01	factor_learning rate = 0.0001
Epoch 42:	error 1883.49	factor_learning rate = 0.0002
Epoch 43:	error 1883.26	factor_learning rate = 0.0004
Epoch 44:	error 1886.21	factor_learning rate = 0.0001
Epoch 45:	error 1882.58	factor_learning rate = 0.0002

```

Epoch 46: error 1882.29      factor_learning_rate = 0.0004
Epoch 47: error 1883.97      factor_learning_rate = 0.0001
Epoch 48: error 1881.58      factor_learning_rate = 0.0002
Epoch 49: error 1881.37      factor_learning_rate = 0.0004
Epoch 50: error 1883.04      factor_learning_rate = 0.0001
Epoch 51: error 1880.74      factor_learning_rate = 0.0002
Epoch 52: error 1880.64      factor_learning_rate = 0.0004
Epoch 53: error 1884.17      factor_learning_rate = 0.0001
Epoch 54: error 1879.83      factor_learning_rate = 0.0002
Epoch 55: error 1879.67      factor_learning_rate = 0.0004
Epoch 56: error 1881.30      factor_learning_rate = 0.0001
Epoch 57: error 1878.88      factor_learning_rate = 0.0002
Epoch 58: error 1878.64      factor_learning_rate = 0.0004

```

```

val_err = squared_error(validation, res=res)
print('Initial solutions as all-1s vectors')
print(f'Final training error = {res.err[-1]:5.3f}')
    f'\tRelative error = {res.err[-1]/squared_error(data)}'
print(f'Validation error = {val_err:5.3f}'
    f'\tRelative error = {val_err/squared_error(validation)}')

```

```

Initial solutions as all-1s vectors
Final training error = 1868.113 Relative error = 0.03338051972473484
Validation error = 5438.536 Relative error = 14.263972215757216

```

#The validation error is far higher than in other approaches

```

#b) Factor matrices and bias vectors are all-0s matrices and vectors. (No normalization.)
# Initial factors are given as a tuple (L, R); R is transposed
init_LR = (np.zeros(shape=(n,10)), np.zeros(shape=(m,10)))
# Initial biases are a tuple (global_bias, row_bias, col_bias)
# Give the true global bias.
init_biases = (avg, np.zeros(shape=n), np.zeros(shape=m))

# You can change the learning rates and bold driver factor if you want
res = sgd(data, 10, bold_driver=2, factor_learning_rate=0.0001, bias_learning_rate=0.0001, ve

```

```

Epoch 41: error NAN      factor_learning_rate = 0.0001
Epoch 42: error NAN      factor_learning_rate = 0.0001
Epoch 43: error NAN      factor_learning_rate = 0.0001
Epoch 44: error NAN      factor_learning_rate = 0.0001
Epoch 45: error NAN      factor_learning_rate = 0.0001
Epoch 46: error NAN      factor_learning_rate = 0.0001
Epoch 47: error NAN      factor_learning_rate = 0.0001
Epoch 48: error NAN      factor_learning_rate = 0.0001
Epoch 49: error NAN      factor_learning_rate = 0.0001
Epoch 50: error NAN      factor_learning_rate = 0.0001
Epoch 51: error NAN      factor_learning_rate = 0.0001
Epoch 52: error NAN      factor_learning_rate = 0.0001
Epoch 53: error NAN      factor_learning_rate = 0.0001
Epoch 54: error NAN      factor_learning_rate = 0.0001
Epoch 55: error NAN      factor_learning_rate = 0.0001
Epoch 56: error NAN      factor_learning_rate = 0.0001
Epoch 57: error NAN      factor_learning_rate = 0.0001

```

```

Epoch  57: error NAN      factor_learning rate = 0.0001
Epoch  58: error NAN      factor_learning rate = 0.0001
Epoch  59: error NAN      factor_learning rate = 0.0001
Epoch  60: error NAN      factor_learning rate = 0.0001
Epoch  61: error NAN      factor_learning rate = 0.0001
Epoch  62: error NAN      factor_learning rate = 0.0001
Epoch  63: error NAN      factor_learning rate = 0.0001
Epoch  64: error NAN      factor_learning rate = 0.0001
Epoch  65: error NAN      factor_learning rate = 0.0001
Epoch  66: error NAN      factor_learning rate = 0.0001
Epoch  67: error NAN      factor_learning rate = 0.0001

Epoch  68: error NAN      factor_learning rate = 0.0001
Epoch  69: error NAN      factor_learning rate = 0.0001
Epoch  70: error NAN      factor_learning rate = 0.0001
Epoch  71: error NAN      factor_learning rate = 0.0001
Epoch  72: error NAN      factor_learning rate = 0.0001
Epoch  73: error NAN      factor_learning rate = 0.0001
Epoch  74: error NAN      factor_learning rate = 0.0001
Epoch  75: error NAN      factor_learning rate = 0.0001
Epoch  76: error NAN      factor_learning rate = 0.0001
Epoch  77: error NAN      factor_learning rate = 0.0001
Epoch  78: error NAN      factor_learning rate = 0.0001
Epoch  79: error NAN      factor_learning rate = 0.0001
Epoch  80: error NAN      factor_learning rate = 0.0001
Epoch  81: error NAN      factor_learning rate = 0.0001
Epoch  82: error NAN      factor_learning rate = 0.0001
Epoch  83: error NAN      factor_learning rate = 0.0001
Epoch  84: error NAN      factor_learning rate = 0.0001
Epoch  85: error NAN      factor_learning rate = 0.0001
Epoch  86: error NAN      factor_learning rate = 0.0001
Epoch  87: error NAN      factor_learning rate = 0.0001
Epoch  88: error NAN      factor_learning rate = 0.0001
Epoch  89: error NAN      factor_learning rate = 0.0001
Epoch  90: error NAN      factor_learning rate = 0.0001
Epoch  91: error NAN      factor_learning rate = 0.0001
Epoch  92: error NAN      factor_learning rate = 0.0001
Epoch  93: error NAN      factor_learning rate = 0.0001
Epoch  94: error NAN      factor_learning rate = 0.0001
Epoch  95: error NAN      factor_learning rate = 0.0001
Epoch  96: error NAN      factor_learning rate = 0.0001
Epoch  97: error NAN      factor_learning rate = 0.0001
Epoch  98: error NAN      factor_learning rate = 0.0001
Epoch  99: error NAN      factor_learning rate = 0.0001

```

```

val_err = squared_error(validation, res=res)
print('Initial solutions as all-1s vectors')
print(f'Final training error = {res.err[-1]:5.3f}')
    f'\tRelative error = {res.err[-1]/squared_error(data)}'
print(f'    Validation error = {val_err:5.3f}')
    f'\tRelative error = {val_err/squared_error(validation)}'

Initial solutions as all-1s vectors
Final training error =    nan    Relative error = nan
    Validation error =    nan    Relative error = nan

```

```
#Thus we see that in the zero case we do not have any solution
```

```
#np.random.normal(0, 1, (3, 3))
# Initial factors are given as a tuple (L, R); R is transposed
init_LR = (np.random.normal(0, 1, (n, 10)), np.random.normal(0, 1, (m, 10)))
# Initial biases are a tuple (global_bias, row_bias, col_bias)
# Give the true global bias.
#init_biases = (avg, np.random.normal(0, 1, (n, 10)), np.random.normal(0, 1, (m, 10)))
```

```
init_LR
```

```
(array([[-1.17425096,  0.19918981, -0.60222315, ...,  0.94614608,
       0.72020795, -0.18975039],
       [ 0.12938449,  2.18551624,  1.17862941, ..., -1.60806639,
      -0.12300904, -0.36229804],
       [-1.12518146,  0.20165542, -0.70840878, ..., -0.99945708,
       0.18227732,  0.28327222],
       ...,
       [-0.95893046, -0.5124027 , -0.95621689, ...,  0.33078595,
       0.21427449, -0.35070937],
       [ 1.05471018, -1.59608667,  1.17445045, ...,  0.33615226,
       2.22510379, -0.94514869],
       [-0.5575909 ,  0.30100682, -2.3460474 , ...,  0.22825616,
      -1.41118481, -1.94459559]]),
array([[-1.05958836, -0.22664738, -0.96912525, ..., -2.11702163,
       -1.40406968, -1.03730845],
       [-0.97952757,  1.242295 , -0.54230601, ..., -0.63007592,
       0.15228372,  0.97380255],
       [-0.33820309,  0.31873166,  0.1123456 , ..., -0.30794351,
       -0.75334118, -1.86813542],
       ...,
       [ 0.44960496,  0.24646594,  0.66482954, ..., -0.21475797,
      -0.69779164,  0.72185682],
       [-0.36235038,  0.39441899, -0.51329037, ...,  0.57037241,
      -0.70571621, -0.28587571],
       [ 0.00957147, -0.37847174,  0.22947165, ..., -0.30678886,
       1.90335051,  1.95844878]]))
```

```
res = sgd(data, 10, bold_driver=2, factor_learning_rate=0.0001, bias_learning_rate=0.0001, ve
```

```
Epoch  41: error NAN    factor_learning rate = 0.0001
Epoch  42: error NAN    factor_learning rate = 0.0001
Epoch  43: error NAN    factor_learning rate = 0.0001
Epoch  44: error NAN    factor_learning rate = 0.0001
Epoch  45: error NAN    factor_learning rate = 0.0001
Epoch  46: error NAN    factor_learning rate = 0.0001
Epoch  47: error NAN    factor_learning rate = 0.0001
Epoch  48: error NAN    factor_learning rate = 0.0001
```

```
Epoch 49: error NAN      factor_learning rate = 0.0001
Epoch 50: error NAN      factor_learning rate = 0.0001
Epoch 51: error NAN      factor_learning rate = 0.0001
Epoch 52: error NAN      factor_learning rate = 0.0001
Epoch 53: error NAN      factor_learning rate = 0.0001
Epoch 54: error NAN      factor_learning rate = 0.0001
Epoch 55: error NAN      factor_learning rate = 0.0001
Epoch 56: error NAN      factor_learning rate = 0.0001
Epoch 57: error NAN      factor_learning rate = 0.0001
Epoch 58: error NAN      factor_learning rate = 0.0001
Epoch 59: error NAN      factor_learning rate = 0.0001

Epoch 60: error NAN      factor_learning rate = 0.0001
Epoch 61: error NAN      factor_learning rate = 0.0001
Epoch 62: error NAN      factor_learning rate = 0.0001
Epoch 63: error NAN      factor_learning rate = 0.0001
Epoch 64: error NAN      factor_learning rate = 0.0001
Epoch 65: error NAN      factor_learning rate = 0.0001
Epoch 66: error NAN      factor_learning rate = 0.0001
Epoch 67: error NAN      factor_learning rate = 0.0001
Epoch 68: error NAN      factor_learning rate = 0.0001
Epoch 69: error NAN      factor_learning rate = 0.0001
Epoch 70: error NAN      factor_learning rate = 0.0001
Epoch 71: error NAN      factor_learning rate = 0.0001
Epoch 72: error NAN      factor_learning rate = 0.0001
Epoch 73: error NAN      factor_learning rate = 0.0001
Epoch 74: error NAN      factor_learning rate = 0.0001
Epoch 75: error NAN      factor_learning rate = 0.0001
Epoch 76: error NAN      factor_learning rate = 0.0001
Epoch 77: error NAN      factor_learning rate = 0.0001
Epoch 78: error NAN      factor_learning rate = 0.0001
Epoch 79: error NAN      factor_learning rate = 0.0001
Epoch 80: error NAN      factor_learning rate = 0.0001
Epoch 81: error NAN      factor_learning rate = 0.0001
Epoch 82: error NAN      factor_learning rate = 0.0001
Epoch 83: error NAN      factor_learning rate = 0.0001
Epoch 84: error NAN      factor_learning rate = 0.0001
Epoch 85: error NAN      factor_learning rate = 0.0001
Epoch 86: error NAN      factor_learning rate = 0.0001
Epoch 87: error NAN      factor_learning rate = 0.0001
Epoch 88: error NAN      factor_learning rate = 0.0001
Epoch 89: error NAN      factor_learning rate = 0.0001
Epoch 90: error NAN      factor_learning rate = 0.0001
Epoch 91: error NAN      factor_learning rate = 0.0001
Epoch 92: error NAN      factor_learning rate = 0.0001
Epoch 93: error NAN      factor_learning rate = 0.0001
Epoch 94: error NAN      factor_learning rate = 0.0001
Epoch 95: error NAN      factor_learning rate = 0.0001
Epoch 96: error NAN      factor_learning rate = 0.0001
Epoch 97: error NAN      factor_learning rate = 0.0001
Epoch 98: error NAN      factor_learning rate = 0.0001
Epoch 99: error NAN      factor_learning rate = 0.0001
```

#So we see that although we are able to generate L R initial solution matrices, the final
solution is a not a number

```
#Factor matrices are random integer matrices with values from 0 to 6, normalized.
L= np.random.randint(7, size=(n,10)) #Create a random matrix of size n by 10 and having value
R = np.random.randint(7, size=(m,10))
print("Original Matrix:")
print(L,R)
Lmax, Lmin = L.max(), L.min()
Rmax, Rmin = R.max(), R.min()
L = (L - Lmin)/(Lmax - Lmin)
R = (R - Rmin) / (Rmax - Rmin)
print("After normalization:")
print(L,R)

Original Matrix:
[[6 2 5 ... 3 3 5]
 [0 3 4 ... 2 0 1]
 [6 3 1 ... 3 0 4]
 ...
 [5 2 3 ... 2 4 0]
 [1 3 0 ... 5 1 0]
 [6 2 5 ... 0 3 2]] [[2 5 4 ... 0 1 1]
 [5 0 4 ... 3 1 0]
 [0 4 1 ... 3 3 6]
 ...
 [5 6 4 ... 6 0 3]
 [3 2 5 ... 1 1 4]
 [4 4 6 ... 0 3 5]]
After normalization:
[[1.          0.33333333 0.83333333 ... 0.5       0.5       0.83333333]
 [0.          0.5         0.66666667 ... 0.33333333 0.          0.16666667]
 [1.          0.5         0.16666667 ... 0.5       0.          0.66666667]
 ...
 [0.83333333 0.33333333 0.5         ... 0.33333333 0.66666667 0.          ]
 [0.16666667 0.5         0.          ... 0.83333333 0.16666667 0.          ]
 [1.          0.33333333 0.83333333 ... 0.          0.5         0.33333333] [[0.33333333
 [0.83333333 0.          0.66666667 ... 0.5       0.16666667 0.          ]
 [0.          0.66666667 0.16666667 ... 0.5       0.5         1.          ]
 ...
 [0.83333333 1.          0.66666667 ... 1.          0.          0.5         ]
 [0.5         0.33333333 0.83333333 ... 0.16666667 0.16666667 0.66666667]
 [0.66666667 0.66666667 1.          ... 0.          0.5         0.83333333]]
```

Initial factors are given as a tuple (L, R); R is transposed

```
init_LR = (L, R)
#
```

```
res = sgd(data, 10, bold_driver=2, factor_learning_rate=0.0001, bias_learning_rate=0.0001, ve
```

```
Epoch 41: error 1280.91      factor_learning_rate = 0.0008
Epoch 42: error 1282.83      factor_learning_rate = 0.0001
Epoch 43: error 1280.21      factor_learning_rate = 0.0002
Epoch 44: error 1278.92      factor_learning_rate = 0.0004
Epoch 45: error 1278.72      factor_learning_rate = 0.0008
```

```

Epoch 46: error 1280.45          factor_learning_rate = 0.0001
Epoch 47: error 1278.12          factor_learning_rate = 0.0002
Epoch 48: error 1276.82          factor_learning_rate = 0.0004
Epoch 49: error 1276.88          factor_learning_rate = 0.0001
Epoch 50: error 1275.91          factor_learning_rate = 0.0002
Epoch 51: error 1275.43          factor_learning_rate = 0.0004
Epoch 52: error 1275.53          factor_learning_rate = 0.0001
Epoch 53: error 1274.84          factor_learning_rate = 0.0002
Epoch 54: error 1274.48          factor_learning_rate = 0.0004
Epoch 55: error 1274.74          factor_learning_rate = 0.0001
Epoch 56: error 1273.97          factor_learning_rate = 0.0002
Epoch 57: error 1273.60          factor_learning_rate = 0.0004
Epoch 58: error 1274.03          factor_learning_rate = 0.0001
Epoch 59: error 1273.11          factor_learning_rate = 0.0002
Epoch 60: error 1272.72          factor_learning_rate = 0.0004

Epoch 61: error 1273.40          factor_learning_rate = 0.0001
Epoch 62: error 1272.32          factor_learning_rate = 0.0002
Epoch 63: error 1271.91          factor_learning_rate = 0.0004
Epoch 64: error 1272.21          factor_learning_rate = 0.0001
Epoch 65: error 1271.43          factor_learning_rate = 0.0002
Epoch 66: error 1271.02          factor_learning_rate = 0.0004
Epoch 67: error 1271.39          factor_learning_rate = 0.0001
Epoch 68: error 1270.61          factor_learning_rate = 0.0002
Epoch 69: error 1270.21          factor_learning_rate = 0.0004
Epoch 70: error 1270.37          factor_learning_rate = 0.0001
Epoch 71: error 1269.76          factor_learning_rate = 0.0002
Epoch 72: error 1269.39          factor_learning_rate = 0.0004
Epoch 73: error 1269.58          factor_learning_rate = 0.0001
Epoch 74: error 1269.00          factor_learning_rate = 0.0002
Epoch 75: error 1268.67          factor_learning_rate = 0.0004
Epoch 76: error 1268.86          factor_learning_rate = 0.0001
Epoch 77: error 1268.21          factor_learning_rate = 0.0002
Epoch 78: error 1267.82          factor_learning_rate = 0.0004
Epoch 79: error 1268.14          factor_learning_rate = 0.0001
Epoch 80: error 1267.45          factor_learning_rate = 0.0002
Epoch 81: error 1267.11          factor_learning_rate = 0.0004
Epoch 82: error 1267.33          factor_learning_rate = 0.0001
Epoch 83: error 1266.69          factor_learning_rate = 0.0002
Epoch 84: error 1266.34          factor_learning_rate = 0.0004
Epoch 85: error 1266.95          factor_learning_rate = 0.0001
Epoch 86: error 1266.01          factor_learning_rate = 0.0002
Epoch 87: error 1265.60          factor_learning_rate = 0.0004
Epoch 88: error 1265.78          factor_learning_rate = 0.0001
Epoch 89: error 1265.19          factor_learning_rate = 0.0002
Epoch 90: error 1264.86          factor_learning_rate = 0.0004
Epoch 91: error 1265.12          factor_learning_rate = 0.0001
Epoch 92: error 1264.45          factor_learning_rate = 0.0002
Epoch 93: error 1264.19          factor_learning_rate = 0.0004
Epoch 94: error 1264.43          factor_learning_rate = 0.0001
Epoch 95: error 1263.78          factor_learning_rate = 0.0002
Epoch 96: error 1263.44          factor_learning_rate = 0.0004
Epoch 97: error 1263.84          factor_learning_rate = 0.0001
Epoch 98: error 1263.07          factor_learning_rate = 0.0002
Epoch 99: error 1262.73          factor_learning_rate = 0.0004

```

```
val_err = squared_error(validation, res=res)
```

https://colab.research.google.com/drive/1RfGKVdaszG6EuUfZlfwf0P4pPSzHUjTK#scrollTo=M3_WAKv4RIRG&printMode=true

```

print('Initial solutions as all-1s vectors')
print(f'Final training error = {res.err[-1]:5.3f}')
    f'\tRelative error = {res.err[-1]/squared_error(data)}')
print(f'Validation error = {val_err:5.3f}'
    f'\tRelative error = {val_err/squared_error(validation)})')

Initial solutions as all-1s vectors
Final training error = 1262.930 Relative error = 0.022566763306068317
Validation error = 289.776 Relative error = 0.7600123874348135

learning_rates = [1e-5,1e-4,1e-3]

start = time.process_time()
err = []
for rate in learning_rates:
    err.append( cv(data, folds=5, k=10, bold_driver=2, factor_learning_rate=rate, bias_learni
print('Rate\terr')
for i in range(len(learning_rates)):
    print(f'{learning_rates[i]}\t{err[i]}')
print(time.process_time() - start)

Fold 1/5
    Error = 2522.5388802632397
Fold 2/5
    Error = 2489.1905454788107
Fold 3/5
    Error = 2504.4814770818143
Fold 4/5
    Error = 2603.2898195405733
Fold 5/5
    Error = 2436.963867852763
Fold 1/5
    Error = 2606.133391562352
Fold 2/5
    Error = 2515.1371980337613
Fold 3/5
    Error = 2425.8187872819076
Fold 4/5
    Error = 2356.2555134840804
Fold 5/5
    Error = 2641.0567237722935
Fold 1/5
    Error = 2543.449346384727
Fold 2/5
    Error = 2723.9636613255407
Fold 3/5
    Error = 2485.382732583863
Fold 4/5
    Error = 2605.006823579612
Fold 5/5
    Error = 2428.441760061842
Rate    err
1e-05   2511.2929180434403

```

```
0.0001 2508.880322826879
0.001 2557.248864787117
2956.489249872
```

```
#The time taken in code above is 2956.48 SECs.The best parameter is for
# learning rate 0.0001 for which the average error is the lowest.
#data = [data]
#k = [10]

#from sklearn.model_selection import GridSearchCV
#param_grid = {'factor_regu':factor_regu, 'bias_regu': bias_regu,'k':k,'factor_learning_rates'

regularization_coeff = [1e-6,1e-5,1e-4]

start = time.process_time()
err = []
for coeff in regularization_coeff:
    err.append( cv(data, folds=5, k=10, bold_driver=2, factor_regu=coeff, factor_learning_rate=0.0001))
print('Rate\terr')
for i in range(len(regularization_coeff)):
    print(f'{regularization_coeff[i]}\t{err[i]}')
print(time.process_time() - start)

Fold 1/5
    Error = 2743.658330049852
Fold 2/5
    Error = 2961.9016783817465
Fold 3/5
    Error = 2748.7784574101943
Fold 4/5
    Error = 2969.309548115035
Fold 5/5
    Error = 2633.911537871388
Fold 1/5
    Error = 2935.336981518091
Fold 2/5
    Error = 2720.21982641188
Fold 3/5
    Error = 2864.97255927242
Fold 4/5
    Error = 2854.3106561156783
Fold 5/5
    Error = 2835.735287572032
Fold 1/5
    Error = 2574.718204069597
Fold 2/5
    Error = 2654.184835630396
Fold 3/5
    Error = 2780.8703271832537
Fold 4/5
    Error = 2764.8582985545318
Fold 5/5
```

```
Error = 2651.355910728593
Rate    err
1e-06  2811.511910365643
1e-05  2842.1150621780203
0.0001 2685.1975152332743
2495.0275192910003
```

#We get the lowest error for REG COEFF 0.0001 . Total time taken was 2495 secs.

#Now doing SGD with regularization L2-Regularization, bold driver of 2, factor learning rate
`res = sgd(data, 10, bold_driver=2, factor_learning_rate=0.0001, bias_learning_rate=0.0001, f`

Epoch 0:	error 628968.60	factor_learning_rate = 0.0001
Epoch 1:	error 270501.48	factor_learning_rate = 0.0002
Epoch 2:	error 122283.66	factor_learning_rate = 0.0004
Epoch 3:	error 64240.54	factor_learning_rate = 0.0008
Epoch 4:	error 37100.29	factor_learning_rate = 0.0016
Epoch 5:	error 21236.32	factor_learning_rate = 0.0032
Epoch 6:	error 11513.69	factor_learning_rate = 0.0064
Epoch 7:	error 6153.06	factor_learning_rate = 0.01
Epoch 8:	error 3830.75	factor_learning_rate = 0.01
Epoch 9:	error 2860.91	factor_learning_rate = 0.01
Epoch 10:	error 2420.47	factor_learning_rate = 0.01
Epoch 11:	error 2190.95	factor_learning_rate = 0.01
Epoch 12:	error 2031.73	factor_learning_rate = 0.01
Epoch 13:	error 1925.51	factor_learning_rate = 0.01
Epoch 14:	error 1833.90	factor_learning_rate = 0.01
Epoch 15:	error 1790.35	factor_learning_rate = 0.01
Epoch 16:	error 1753.30	factor_learning_rate = 0.01
Epoch 17:	error 1692.71	factor_learning_rate = 0.01
Epoch 18:	error 1692.06	factor_learning_rate = 0.01
Epoch 19:	error 1670.95	factor_learning_rate = 0.01
Epoch 20:	error 1622.03	factor_learning_rate = 0.01
Epoch 21:	error 1638.80	factor_learning_rate = 0.0001
Epoch 22:	error 1516.26	factor_learning_rate = 0.0002
Epoch 23:	error 1424.12	factor_learning_rate = 0.0004
Epoch 24:	error 1367.12	factor_learning_rate = 0.0008
Epoch 25:	error 1338.52	factor_learning_rate = 0.0016
Epoch 26:	error 1333.34	factor_learning_rate = 0.0032
Epoch 27:	error 1355.56	factor_learning_rate = 0.0001
Epoch 28:	error 1333.34	factor_learning_rate = 0.0002
Epoch 29:	error 1317.81	factor_learning_rate = 0.0004
Epoch 30:	error 1309.73	factor_learning_rate = 0.0008
Epoch 31:	error 1308.74	factor_learning_rate = 0.0016
Epoch 32:	error 1313.39	factor_learning_rate = 0.0001
Epoch 33:	error 1306.68	factor_learning_rate = 0.0002
Epoch 34:	error 1302.11	factor_learning_rate = 0.0004
Epoch 35:	error 1300.16	factor_learning_rate = 0.0008
Epoch 36:	error 1301.27	factor_learning_rate = 0.0001
Epoch 37:	error 1298.98	factor_learning_rate = 0.0002
Epoch 38:	error 1297.56	factor_learning_rate = 0.0004
Epoch 39:	error 1297.17	factor_learning_rate = 0.0008
Epoch 40:	error 1300.09	factor_learning_rate = 0.0001
Epoch 41:	error 1296.70	factor_learning_rate = 0.0002
Epoch 42:	error 1295.04	factor_learning_rate = 0.0004

```

Epoch 43: error 1295.07      factor_learning_rate = 0.0001
Epoch 44: error 1294.03      factor_learning_rate = 0.0002
Epoch 45: error 1293.50      factor_learning_rate = 0.0004
Epoch 46: error 1293.77      factor_learning_rate = 0.0001
Epoch 47: error 1292.93      factor_learning_rate = 0.0002
Epoch 48: error 1292.46      factor_learning_rate = 0.0004
Epoch 49: error 1292.73      factor_learning_rate = 0.0001
Epoch 50: error 1291.84      factor_learning_rate = 0.0002
Epoch 51: error 1291.45      factor_learning_rate = 0.0004
Epoch 52: error 1291.72      factor_learning_rate = 0.0001
Epoch 53: error 1290.90      factor_learning_rate = 0.0002
Epoch 54: error 1290.47      factor_learning_rate = 0.0004
Epoch 55: error 1290.61      factor_learning_rate = 0.0001
Epoch 56: error 1289.90      factor_learning_rate = 0.0002
Epoch 57: error 1289.52      factor_learning_rate = 0.0004
Epoch 58: error 1289.59      factor_learning_rate = 0.0001

```

```

val_err = squared_error(validation, res=res)
print('Initial solutions as all-1s vectors')
print(f'Final training error = {res.err[-1]:5.3f}')
    f'\tRelative error = {res.err[-1]/squared_error(data)}'
print(f'Validation error = {val_err:5.3f}'
    f'\tRelative error = {val_err/squared_error(validation)}')

```

```

Initial solutions as all-1s vectors
Final training error = 1277.875 Relative error = 0.02283381383939651
Validation error = 233.456  Relative error = 0.6122996982054809

```

#This result is lower than the best we got earlier, because the bold driver factor used there
#So we need to work for that, and see the best parameters corresponding to it.

```

learning_rates = [1e-5, 1e-4, 1e-3]

start = time.process_time()
err = []
for rate in learning_rates:
    err.append(cv(data, folds=5, k=10, bold_driver=3, factor_learning_rate=rate, bias_learni
print('Rate\tterr')
for i in range(len(learning_rates)):
    print(f'{learning_rates[i]}\t{err[i]}')
print(time.process_time() - start)

Fold 1/5
    Error = 2722.5288581181385
Fold 2/5
    Error = 2548.302374176806
Fold 3/5
    Error = 2531.347277367209
Fold 4/5
    Error = 2540.238496009842
Fold 5/5

```

```

        Error = 2748.2896215624846
Fold 1/5
        Error = 2633.391336263878
Fold 2/5
        Error = 2650.69334040362
Fold 3/5
        Error = 2542.0456009362965
Fold 4/5
        Error = 2656.0247062022363
Fold 5/5
        Error = 2449.8824910680437
Fold 1/5
        Error = 2312.145161770416
Fold 2/5
        Error = 2552.8925029346537
Fold 3/5
        Error = 2418.1880469407265
Fold 4/5
        Error = 2353.8519240647397
Fold 5/5
        Error = 2659.890620634777
Rate    err
1e-05  2618.141325446896
0.0001 2586.4074949748147
0.001   2459.3936512690625
2978.851456674999

```

#The error is minimum for 0.001 value. So I use this to get the next parameter optimum value

```
regularization_coeff = [1e-6, 1e-5, 1e-4]
```

```

start = time.process_time()
err = []
for coeff in regularization_coeff:
    err.append(cv(data, folds=5, k=10, bold_driver=3, factor_regu=coeff, factor_learning_rate=0.001))
print('Rate\terr')
for i in range(len(regularization_coeff)):
    print(f'{regularization_coeff[i]}\t{err[i]}')
print(time.process_time() - start)

```

```

Fold 1/5
        Error = 2801.487380575274
Fold 2/5
        Error = 2820.085083191071
Fold 3/5
        Error = 2947.297234022993
Fold 4/5
        Error = 3017.180038227253
Fold 5/5
        Error = 3011.9426937494577
Fold 1/5
        Error = 2956.9260771417444
Fold 2/5
        Error = 2868.551920255637

```

```

Fold 3/5
    Error = 2764.358354477345
Fold 4/5
    Error = 3131.893527184753
Fold 5/5
    Error = 2957.233345347008
Fold 1/5
    Error = 2743.44515320909
Fold 2/5
    Error = 2595.991925187045
Fold 3/5
    Error = 2684.655141344161
Fold 4/5
    Error = 2688.8406079075107
Fold 5/5
    Error = 2750.029991512247
Rate   err
1e-06  2919.5984859532095
1e-05  2935.7926448812973
0.0001 2692.5925638320105
2531.3175626440006

```

```
#So it took 2531 seconds to complete and the lowest error came for 0.0001 as regularization c
```

```
#Now doing SGD with regularization L2-Regularization, bold driver of 3, factor learning rate
res = sgd(data, 10, bold_driver=3, factor_learning_rate=0.001, bias_learning_rate=0.001, fact
```

epochn	2/ epoch error 1291./4	factor_learning rate = 0.003
Epoch	28: error 1322.85	factor_learning rate = 0.001
Epoch	29: error 1286.62	factor_learning rate = 0.003
Epoch	30: error 1317.55	factor_learning rate = 0.001
Epoch	31: error 1283.19	factor_learning rate = 0.003
Epoch	32: error 1323.15	factor_learning rate = 0.001
Epoch	33: error 1278.04	factor_learning rate = 0.003
Epoch	34: error 1315.06	factor_learning rate = 0.001
Epoch	35: error 1274.85	factor_learning rate = 0.003
Epoch	36: error 1309.03	factor_learning rate = 0.001
Epoch	37: error 1272.15	factor_learning rate = 0.003
Epoch	38: error 1305.50	factor_learning rate = 0.001
Epoch	39: error 1269.99	factor_learning rate = 0.003
Epoch	40: error 1302.88	factor_learning rate = 0.001
Epoch	41: error 1265.51	factor_learning rate = 0.003
Epoch	42: error 1295.56	factor_learning rate = 0.001
Epoch	43: error 1262.78	factor_learning rate = 0.003
Epoch	44: error 1294.25	factor_learning rate = 0.001
Epoch	45: error 1259.57	factor_learning rate = 0.003
Epoch	46: error 1293.88	factor_learning rate = 0.001
Epoch	47: error 1257.07	factor_learning rate = 0.003
Epoch	48: error 1292.61	factor_learning rate = 0.001
Epoch	49: error 1254.53	factor_learning rate = 0.003
Epoch	50: error 1287.86	factor_learning rate = 0.001
Epoch	51: error 1251.21	factor_learning rate = 0.003
Epoch	52: error 1286.20	factor_learning rate = 0.001
Epoch	53: error 1249.37	factor_learning rate = 0.003
Epoch	54: error 1277.64	factor_learning rate = 0.001
Epoch	55: error 1246.74	factor learning rate = 0.003

```

Epoch  56: error 1280.83      factor_learning_rate = 0.001
Epoch  57: error 1245.58      factor_learning_rate = 0.003
Epoch  58: error 1281.63      factor_learning_rate = 0.001
Epoch  59: error 1241.89      factor_learning_rate = 0.003
Epoch  60: error 1277.92      factor_learning_rate = 0.001
Epoch  61: error 1239.59      factor_learning_rate = 0.003
Epoch  62: error 1277.06      factor_learning_rate = 0.001
Epoch  63: error 1238.46      factor_learning_rate = 0.003
Epoch  64: error 1272.13      factor_learning_rate = 0.001
Epoch  65: error 1236.15      factor_learning_rate = 0.003
Epoch  66: error 1269.85      factor_learning_rate = 0.001
Epoch  67: error 1233.26      factor_learning_rate = 0.003
Epoch  68: error 1265.91      factor_learning_rate = 0.001
Epoch  69: error 1232.96      factor_learning_rate = 0.003
Epoch  70: error 1265.74      factor_learning_rate = 0.001
Epoch  71: error 1229.62      factor_learning_rate = 0.003
Epoch  72: error 1262.56      factor_learning_rate = 0.001
Epoch  73: error 1227.98      factor_learning_rate = 0.003
Epoch  74: error 1266.34      factor_learning_rate = 0.001
Epoch  75: error 1226.77      factor_learning_rate = 0.003
Epoch  76: error 1261.01      factor_learning_rate = 0.001
Epoch  77: error 1223.48      factor_learning_rate = 0.003
Epoch  78: error 1257.90      factor_learning_rate = 0.001
Epoch  79: error 1224.11      factor_learning_rate = 0.003
Epoch  80: error 1257.64      factor_learning_rate = 0.001
Epoch  81: error 1221.53      factor_learning_rate = 0.003
Epoch  82: error 1255.34      factor_learning_rate = 0.001
Epoch  83: error 1218.77      factor_learning_rate = 0.003
Epoch  84: error 1253.05      factor_learning_rate = 0.001
Epoch  85: error 1218.98      factor_learning_rate = 0.003

```

```

val_err = squared_error(validation, res=res)
print('Initial solutions as all-1s vectors')
print(f'Final training error = {res.err[-1]:5.3f}')
    f'\tRelative error = {res.err[-1]/squared_error(data)}')
print(f'Validation error = {val_err:5.3f}'
    f'\tRelative error = {val_err/squared_error(validation)})'

```

```

Initial solutions as all-1s vectors
Final training error = 1240.353 Relative error = 0.02216334060456452
Validation error = 240.222  Relative error = 0.6300442533451417

```

#However, this validation score is greater than 188 what we obtained earlier:

```

#res = sgd(data, 10, bold_driver=3, factor_learning_rate=0.01)
#So we try those learning rates as well

```

```
learning_rates = [0.01, 0.1, 1e-5, 1e-4, 1e-3]
```

```

start = time.process_time()
err = []
for rate in learning_rates:

```

```
err.append( cv(data, folds=5, k=10, bold_driver=3, factor_learning_rate=rate, bias_learni
print('Rate\terr')
for i in range(len(learning_rates)):
    print(f'{learning_rates[i]}\t{err[i]}')
print(time.process_time() - start)

Error is 5./2/2b5b5419b0425e+20, you should probably use slower learning rate
Error is 1.3163441052078707e+294, you should probably use slower learning rate
Error is 9.416417444054082e+224, you should probably use slower learning rate
Error is 3.2342781242768822e+227, you should probably use slower learning rate
Error is 3.586297209233935e+124, you should probably use slower learning rate
Error is 1.7155035737386879e+47, you should probably use slower learning rate
Error is 4.55041345993541e+20, you should probably use slower learning rate

Error is 5503149314228485.0, you should probably use slower learning rate
Error is 1827605033038075.5, you should probably use slower learning rate
Error is 1.472385389577297e+299, you should probably use slower learning rate
Error is 734219046760867.5, you should probably use slower learning rate
Error is 1.249741888653658e+164, you should probably use slower learning rate
Error is 322695.57768617227, you should probably use slower learning rate
Error is 2.4200710706305106e+202, you should probably use slower learning rate
Error is 1.101113170840056e+131, you should probably use slower learning rate
Error is 3.623921722199731e+144, you should probably use slower learning rate
Error is 402062.1226076005, you should probably use slower learning rate
Error is 1.8596667068632105e+202, you should probably use slower learning rate
Error is 8.69015602505699e+201, you should probably use slower learning rate
Error is 6.716854006823664e+201, you should probably use slower learning rate
    Error = nan
Fold 1/5
    Error = 2467.853681477283
Fold 2/5
    Error = 2527.732967914114
Fold 3/5
    Error = 2360.838935521543
Fold 4/5
    Error = 2622.725077740915
Fold 5/5
    Error = 2617.765652226571
Fold 1/5
    Error = 2464.9599944767897
Fold 2/5
    Error = 2639.924315673732
Fold 3/5
    Error = 2386.4413731964187
Fold 4/5
    Error = 2471.631361123579
Fold 5/5
    Error = 2555.01243837131
Fold 1/5
    Error = 2621.1776178656646
Fold 2/5
    Error = 2396.846637590643
Fold 3/5
    Error = 2813.4069896910764
Fold 4/5
    Error = 2439.9967283026454
Fold 5/5
    Error = 2362.993120233712
```

```
Rate    err
0.01   2315.183788809335
0.1    nan
1e-05  2519.383262976085
0.0001 2503.5938965683663
0.001   2526.884218736748
5002.484504351001
```

```
#Thus, now we can clearly see that the least error is for step value 0.01
#The total time taken was 5002 seconds
```

```
#Finding the regularizaiton hyper-parameter which is optimal for step value 0.01
regularization_coeff = [1e-6,1e-5,1e-4]

start = time.process_time()
err = []
for coeff in regularization_coeff:
    err.append( cv(data, folds=5, k=10, bold_driver=3, factor_regu=coeff, factor_learning_rate=0.01))
print('Rate\terr')
for i in range(len(regularization_coeff)):
    print(f'{regularization_coeff[i]}\t{err[i]}')
print(time.process_time() - start)

Fold 1/5
    Error = 2723.9522335410074
Fold 2/5
    Error = 2479.8903454863907
Fold 3/5
    Error = 2701.477978900012
Fold 4/5
    Error = 2620.651166682969
Fold 5/5
    Error = 2770.3373583415414
Fold 1/5
    Error = 2896.6547637955846
Fold 2/5
    Error = 2697.6782914859987
Fold 3/5
    Error = 2697.3500426837727
Fold 4/5
    Error = 2666.9147475426907
Fold 5/5
    Error = 2700.9933094694256
Fold 1/5
    Error = 2578.160316807287
Fold 2/5
    Error = 2634.4587197658807
Fold 3/5
    Error = 2378.3291158020434
Fold 4/5
    Error = 2418.9902713865345
Fold 5/5
    Error = 2738.4165664106777
```

Rate	err
1e-06	2659.261816590384
1e-05	2731.918230995495
0.0001	2549.6709980344845
	2285.068527286

```
#So the least error was found for coeff 0.0001 and time taken is 2285 secs
```

```
#Now applying the optimal hyper-parameters
```

```
#Now doing SGD with regularization L2-Regularization, bold driver of 2, factor learning rate
res = sgd(data, 10, bold_driver=3, factor_learning_rate=0.01, bias_learning_rate=0.01, factor
```

epoch	error	factor_learning_rate
42	1407.05	0.01
43	1390.47	0.01
44	1386.31	0.01
45	1366.78	0.01
46	1390.11	0.01
47	1380.69	0.01
48	1406.17	0.01
49	1369.86	0.01
50	1371.19	0.01
51	1386.71	0.01
52	1355.18	0.01
53	1351.27	0.01
54	1362.63	0.01
55	1357.35	0.01
56	1349.00	0.01
57	1311.33	0.01
58	1336.14	0.01
59	1348.36	0.01
60	1346.03	0.01
61	1330.36	0.01
62	1352.84	0.01
63	1321.40	0.01
64	1355.92	0.01
65	1325.09	0.01
66	1331.01	0.01
67	1318.82	0.01
68	1334.22	0.01
69	1307.41	0.01
70	1318.32	0.01
71	1325.00	0.01
72	1314.91	0.01
73	1304.17	0.01
74	1300.56	0.01
75	1293.69	0.01
76	1306.56	0.01
77	1282.80	0.01
78	1295.34	0.01
79	1301.15	0.01
80	1288.67	0.01
81	1284.35	0.01
82	1279.61	0.01
83	1293.37	0.01
84	1279.02	0.01

```

Epoch  85: error 1286.47      factor_learning_rate = 0.01
Epoch  86: error 1265.21      factor_learning_rate = 0.01
Epoch  87: error 1297.84      factor_learning_rate = 0.01
Epoch  88: error 1277.86      factor_learning_rate = 0.01
Epoch  89: error 1279.30      factor_learning_rate = 0.01
Epoch  90: error 1269.20      factor_learning_rate = 0.01
Epoch  91: error 1246.06      factor_learning_rate = 0.01
Epoch  92: error 1247.97      factor_learning_rate = 0.01
Epoch  93: error 1261.16      factor_learning_rate = 0.01
Epoch  94: error 1255.30      factor_learning_rate = 0.01
Epoch  95: error 1266.86      factor_learning_rate = 0.01
Epoch  96: error 1241.27      factor_learning_rate = 0.01
Epoch  97: error 1267.68      factor_learning_rate = 0.01

```

```

Epoch  98: error 1267.82      factor_learning_rate = 0.01
Epoch  99: error 1309.23      factor_learning_rate = 0.01

```

```

val_err = squared_error(validation, res=res)
print('Initial solutions as all-1s vectors')
print(f'Final training error = {res.err[-1]:5.3f}')
    f'\tRelative error = {res.err[-1]/squared_error(data)}'
print(f'Validation error = {val_err:5.3f}'
    f'\tRelative error = {val_err/squared_error(validation)}')

```

```

Initial solutions as all-1s vectors
Final training error = 1287.652 Relative error = 0.02300850059005928
Validation error = 216.934  Relative error = 0.5689668332814921

```

```

#So now we have a validation score which is just above 188 which might be the
# case for over-fitting. The regularization causes the error to increase
# a bit to 216. This seems to be the optimal model.

```

```

#####
## TASK 3
#Task 3: ICA for housing prices
#In this task, we study the applications of independent component analysis to housing price c
#US. The data set housing_prices.csv4 contains the monthly house price index for twenty metro
#areas in the US from January 1987 to June 2014. As is common to time series data, the rows c
#the locations and the columns to the time stamps. The areas are in housing_prices_locations.
#the time stamps are in housing_prices_times.txt.
#Get yourself familiar with the data by studying which locations it covers and by plotting th
#series.
#This data contains missing values. We start by using k-nearest neighbour imputer. You can p
#again. Did it change?
#Compute the ICA of the data. First, compute only top-4 components and plot them. How would y
#interpret the results?
#The US housing bubble made the housing prices climb heavily from around year 2000 until 2006
#they started to decline. By late 2008, the decline had turned into a crisis with significant
#housing prices; the houses would not start to fully recover until 2012. Can you identify the
#independent components? Repeat ICA with higher number of components (say, 8). Can you identi
#events now? Remember that some signals might have their sign changed, so you need to multipl

```

```
#-1 to see the signal the right way (though of course you need to first decide which signals
#More common way to impute values is to replace the missing values with a mean over the signals
#this and re-compute ICA. Did the results change? What does it mean if the results changed/did
#####
```

```
from sklearn.impute import KNNImputer
from sklearn.decomposition import FastICA

# Read data
data = np.genfromtxt('housing_prices.csv', delimiter=',', skip_header=1, missing_values="NA",
                     encoding='utf-8')

# We'd like to have cities as columns and time stamps as rows, so transpose
#data = data.T

# Read locations
loc=[]
with open('housing_prices_locations.txt', 'r') as f:
    for row in f:
        loc.append(row.strip())

# Read times
times=[]
with open('housing_prices_times.txt', 'r') as f:
    for row in f:
        times.append(row.strip())

# Plot data
for i in range(data.shape[0]):
    plt.subplot(4, 5, i+1)
    plt.title(loc[i])
    plt.plot(data[i,:])
plt.tight_layout()
plt.show()
```



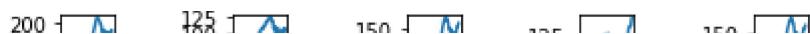
```
# Remove missing values using k-nearest neighbour imputer
```

```
imputer = KNNImputer(n_neighbors=2)
```

```
imp_data=imputer.fit_transform(data)
```



```
#So now we see how the data is imputed
```



```
data
```

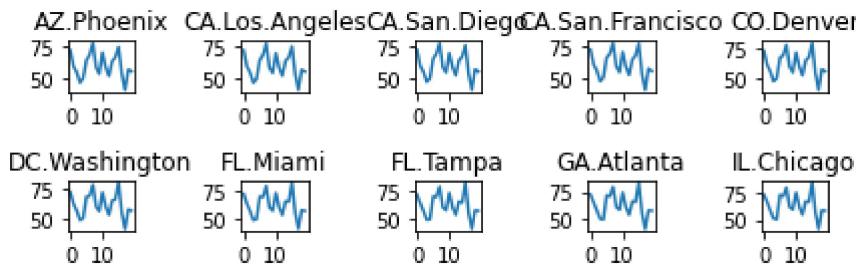
```
array([[  nan,  nan,  nan, ..., 145.4 , 146.05, 146.9 ],
       [ 59.33, 59.65, 59.99, ..., 219.63, 221.94, 223.33],
       [ 54.67, 54.89, 55.16, ..., 200.72, 201.94, 203.32],
       ...,
       [ 41.05, 41.28, 41.06, ..., 165.2 , 167.09, 168.97],
       [  nan,  nan,  nan, ..., 136.73, 138.54, 140.12],
       [  nan,  nan,  nan, ..., 165.74, 168.1 , 169.96]])
```

```
imp_data
```

```
array([[ 71.845, 72.48 , 72.55 , ..., 145.4 , 146.05 , 146.9 ],
       [ 59.33 , 59.65 , 59.99 , ..., 219.63 , 221.94 , 223.33 ],
       [ 54.67 , 54.89 , 55.16 , ..., 200.72 , 201.94 , 203.32 ],
       ...,
       [ 41.05 , 41.28 , 41.06 , ..., 165.2 , 167.09 , 168.97 ],
       [ 56.795, 56.95 , 57.16 , ..., 136.73 , 138.54 , 140.12 ],
       [ 55.545, 55.68 , 55.53 , ..., 165.74 , 168.1 , 169.96 ]])
```

```
# Plot imputed data
```

```
for i in range(imp_data.shape[0]):
    plt.subplot(4, 5, i+1)
    plt.title(loc[i])
    plt.plot(imp_data[:,i])
plt.tight_layout()
plt.show()
```



#We can clearly see that the plots have changed significantly.

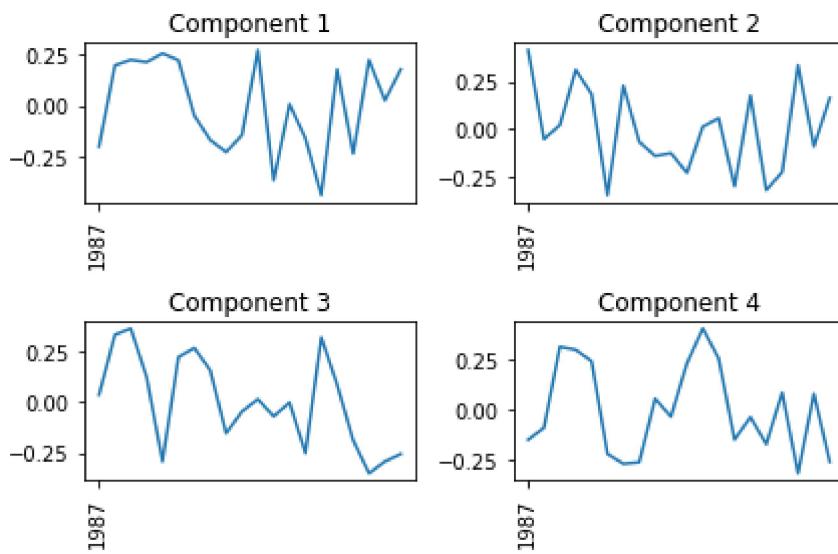
o o o o o

Compute ICA

```
ica = FastICA(n_components=4)
all_comps = ica.fit_transform(imp_data)
```

Plot components

```
for i in range(all_comps.shape[1]):
    ax = plt.subplot(2, 2, i+1)
    ax.set_xticks(np.arange(0, len(times), 36))
    ax.set_xticklabels(np.arange(1987, 2015, 3), rotation='vertical')
    plt.title(f'Component {i+1}')
    plt.plot(all_comps[:, i])
plt.tight_layout()
plt.show()
```



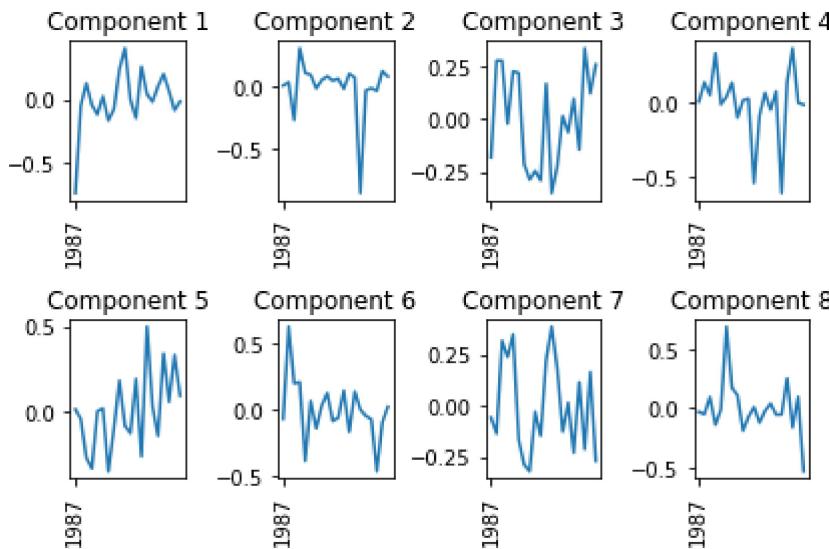
#From the data we can see that the cities can be classified into 4 different categories

In the Component 1 type of cities we can see that the prices of houses show great abberation
zig zag pattern . The rise in price of houses in each of the component can be seen from the
2000 till 2006 and then sudden decline in prices since 2008 is observed in all components,
particularly in component 3 . Compoenent 1 houses still has their price rise in most parts
indicating the cities where investments can be highly profitable.

```
# Compute ICA 8 components
ica = FastICA(n_components=8)
all_comps = ica.fit_transform(imp_data)

/usr/local/lib/python3.7/dist-packages/sklearn/decomposition/_fastica.py:119: ConvergenceWarning:
  ConvergenceWarning)
```

```
# Plot 8 components
for i in range(all_comps.shape[1]):
    ax = plt.subplot(2, 4, i+1)
    ax.set_xticks(np.arange(0, len(times), 36))
    ax.set_xticklabels(np.arange(1987, 2015, 3), rotation='vertical')
    plt.title(f'Component {i+1}')
    plt.plot(all_comps[:, i])
plt.tight_layout()
plt.show()
```

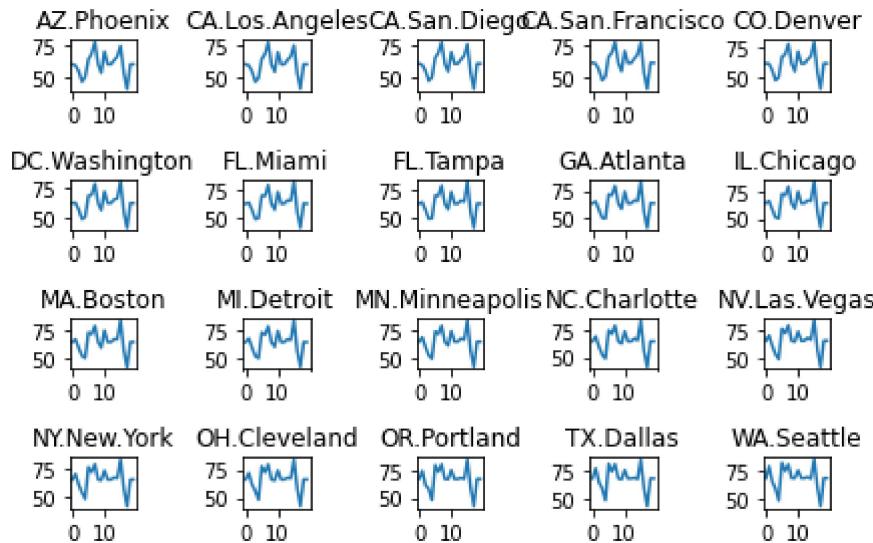


```
#Now we can see more components where the rise in price from 2000 to 2006 can be
#observed and then subsequent crisis fall from 2008 onwards. Component 2 and 4 needs
#to have a flip in sign by multiplying by -1 in order to observe the results in
#the expected way.
```

```
# To impute mean of columns, we use SimpleImputer
from sklearn.impute import SimpleImputer
mean_imputer = SimpleImputer(missing_values=np.nan, strategy='mean')
mean_imp_data = mean_imputer.fit_transform(data)
```

```
# Plot imputed data
for i in range(mean_imp_data.shape[0]):
    plt.subplot(4, 5, i+1)
    plt.title(loc[i])
    plt.plot(mean_imp_data[:, i])
```

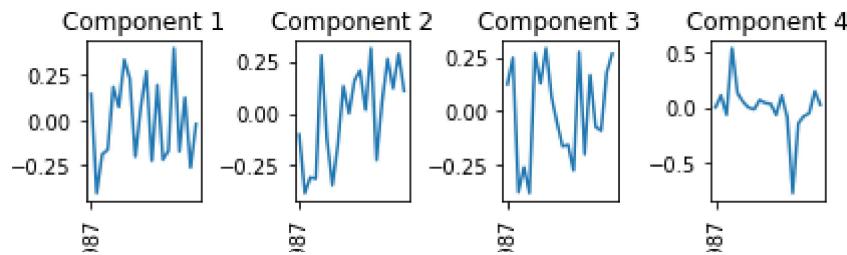
```
plt.tight_layout()
plt.show()
```



```
# Compute ICA 8 components
ica = FastICA(n_components=8)
all_comps = ica.fit_transform(mean_imp_data)
```

```
# Plot imputed data
for i in range(mean_imp_data.shape[0]):
    plt.subplot(4, 5, i+1)
    plt.title(loc[i])
    plt.plot(mean_imp_data[:,i])
plt.tight_layout()
plt.show()
```

```
# Plot 8 components
for i in range(all_comps.shape[1]):
    ax = plt.subplot(2, 4, i+1)
    ax.set_xticks(np.arange(0, len(times), 36))
    ax.set_xticklabels(np.arange(1987, 2015, 3), rotation='vertical')
    plt.title(f'Component {i+1}')
    plt.plot(all_comps[:,i])
plt.tight_layout()
plt.show()
```



We see similar observation even with mean imputation, although, the drop in price is more prominent when we use k-nn. The ICA can help us see which components to invest in.

0.0

```

tmp = np.genfromtxt('ratings.csv', delimiter=',', dtype='int32')
val = np.array(tmp[:,2]+1, dtype='float')/6
# ratings are from 0 to 5, but we can't have 0 entries, as they'd be confused
# with missing ratings.
data = coo_matrix( (val, (tmp[:,0]-1, tmp[:,1]-1)) )
# data is 1-indexed, so we need to reduce 1 to get it 0-indexed

tmp = np.genfromtxt('ratings_validation.csv', delimiter=',', dtype='int32')
val = np.array(tmp[:,2]+1, dtype='float')/6
validation = coo_matrix( (val, (tmp[:,0]-1, tmp[:,1]-1)) )

n = data.get_shape()[0]
m = data.get_shape()[1]
avg = data.sum()/data.count_nonzero()
print(f'Data has {n} rows and {m} columns')
print(f'Average rating is {avg:.4F}')

# Read movies
movies = []
with open('movies.txt') as f:
    for row in f:
        movies.append(row.strip())

# Compute just factors
res = sgd(data, 10, factor_learning_rate=0.001)
val_err = squared_error(validation, res.L, res.R)

```

```

print('Only factors, fixed rate, no regularization')
print(f'Final training error = {res.err[-1]:5.3f}')
    f'\tRelative error = {res.err[-1]/squared_error(data)}'
print(f'Validation error = {val_err:5.3f}'
    f'\tRelative error = {val_err/squared_error(validation)})')

# Plot errors over epochs
plt.plot(res.err)
plt.gca().set_yscale('log') # log scale makes these plots easier to read
plt.show()
plt.close()

# Print the top-10 and bottom-10 movies
print_factors(res.R, movies)

# Compute just factors, but use bold driver and slower learning
res = sgd(data, 10, bold_driver=2, factor_learning_rate=0.0001)
val_err = squared_error(validation, res.L, res.R)
print('Only factors, bold driver, no regularization')
print(f'Final training error = {res.err[-1]:5.3f}'
    f'\tRelative error = {res.err[-1]/squared_error(data)}')
print(f'Validation error = {val_err:5.3f}'
    f'\tRelative error = {val_err/squared_error(validation)})')

# Plot errors over epochs
plt.plot(res.err)
plt.gca().set_yscale('log') # log scale makes these plots easier to read
plt.show()
plt.close()

#####
## YOUR PART STARTS HERE
## Play around with different learning rates & bold driver multipliers
#####

#####
## YOUR PART ENDS HERE
#####

# Compute only bias; factor_learning_rate=None means don't compute factors
res = sgd(data, 10, bold_driver=2, factor_learning_rate=None, bias_learning_rate=0.0001, vert
val_err = squared_error(validation, res.L, res.R, res.global_bias, res.row_bias, res.col_bias
print('Bias only, no regularization')
print(f'Final training error = {res.err[-1]:5.3f}'
    f'\tRelative error = {res.err[-1]/squared_error(data)}')
print(f'Validation error = {val_err:5.3f}'
    f'\tRelative error = {val_err/squared_error(validation)})')

# Compute factors and bias
res = sgd(data, 10, bold_driver=2, factor_learning_rate=0.0001, bias_learning_rate=0.0001, ve

```

```

val_err = squared_error(validation, res.L, res.R, res.global_bias, res.row_bias, res.col_bias)
print('Factors and bias, no regularization')
print(f'Final training error = {res.err[-1]:5.3f}')
    f'\tRelative error = {res.err[-1]/squared_error(data)}'
print(f'Validation error = {val_err:5.3f}'
    f'\tRelative error = {val_err/squared_error(validation)}')

# Print the movies based on their bias
print_factors(res.col_bias, movies)

# Add regularization
res = sgd(data, 10, bold_driver=2, factor_learning_rate=0.0001, bias_learning_rate=0.0001, f
val_err = squared_error(validation, res.L, res.R, res.global_bias, res.row_bias, res.col_bias)
print('Factors and bias, regularized')
print(f'Final training error = {res.err[-1]:5.3f}'
    f'\tRelative error = {res.err[-1]/squared_error(data)}')
print(f'Validation error = {val_err:5.3f}'
    f'\tRelative error = {val_err/squared_error(validation)}')

# To be able to compute the SVD, you can convert the data into full array with
#data.toarray()
# To use squared_error(), you must multiply Sigma to U or V and use a transpose of V.
#####
## YOUR PART STARTS HERE
## Use the parameters from earlier and compare your errors (training and validation)
## to those you get with rank-10 truncated SVD
#####

#####
## YOUR PART ENDS HERE
#####

#####
## TASK 2 ##
#####

# Initial factors are given as a tuple (L, R); R is transposed
init_LR = (np.ones(shape=(n,10)), np.ones(shape=(m,10)))
# Initial biases are a tuple (global_bias, row_bias, col_bias
# Give the true global bias.
init_biases = (avg, np.ones(shape=n), np.ones(shape=m))

# You can change the learning rates and bold driver factor if you want
res = sgd(data, 10, bold_driver=2, factor_learning_rate=0.0001, bias_learning_rate=0.0001, ve
val_err = squared_error(validation, res=res)
print('Initial solutions as all-1s vectors')
print(f'Final training error = {res.err[-1]:5.3f}'
    f'\tRelative error = {res.err[-1]/squared_error(data)}')
print(f'Validation error = {val_err:5.3f}'
    f'\tRelative error = {val_err/squared_error(validation)}')

#####

```

```
## YOUR PART STARTS HERE
## Try the other initial solutions
#####

#####
## YOUR PART ENDS HERE
#####

# try two different values for learning rates using 5-fold cross validation
learning_rates = [1e-5, 1e-3]
err = []
for rate in learning_rates:
    err.append( cv(data, folds=5, k=10, bold_driver=2, factor_learning_rate=rate, bias_learni
print('Rate\tterr')
for i in range(len(learning_rates)):
    print(f'{learning_rates[i]}\t{err[i]}')

#####
## YOUR PART STARTS HERE
## Implement the grid search. After you've found the optimum combination of
## hyperparameters, run sgd with full training data using those parameters and
## compute the validation error.
#####

#####
## YOUR PART ENDS HERE
#####

#####
## TASK 3 ##
#####

from sklearn.impute import KNNImputer
from sklearn.decomposition import FastICA

# Read data
data = np.genfromtxt('housing_prices.csv', delimiter=',', skip_header=1, missing_values="NA",
# Read locations
loc=[]
with open('housing_prices_locations.txt', 'r') as f:
    for row in f:
        loc.append(row.strip())
# Read times
times=[]
with open('housing_prices_times.txt', 'r') as f:
    for row in f:
        times.append(row.strip())
```

```
# Plot data
for i in range(data.shape[0]):
    plt.subplot(4, 5, i+1)
    plt.title(loc[i])
    plt.plot(data[i,:])
plt.tight_layout()
plt.show()

# We'd like to have cities as columns and time stamps as rows, so transpose
data = data.T

# Remove missing values using k-nearest neighbour imputer
imputer = KNNImputer(n_neighbors=2)
imp_data=imputer.fit_transform(data)

#Plot data again

# Compute ICA
ica = FastICA(n_components=4)
all_comps = ica.fit_transform(imp_data)

# Plot components
for i in range(all_comps.shape[1]):
    ax = plt.subplot(2, 2, i+1)
    ax.set_xticks(np.arange(0,len(times),36))
    ax.set_xticklabels(np.arange(1987, 2015, 3), rotation='vertical')
    plt.title(f'Component {i+1}')
    plt.plot(all_comps[:,i])
plt.tight_layout()
plt.show()

# To impute mean of columns, we use SimpleImputer
from sklearn.impute import SimpleImputer
mean_imputer = SimpleImputer(missing_values=np.nan, strategy='mean')
mean_imp_data = mean_imputer.fit_transform(data)

#####
## YOUR PART STARTS HERE
## Do further analysis
#####

#####
## YOUR PART ENDS HERE
#####
```

